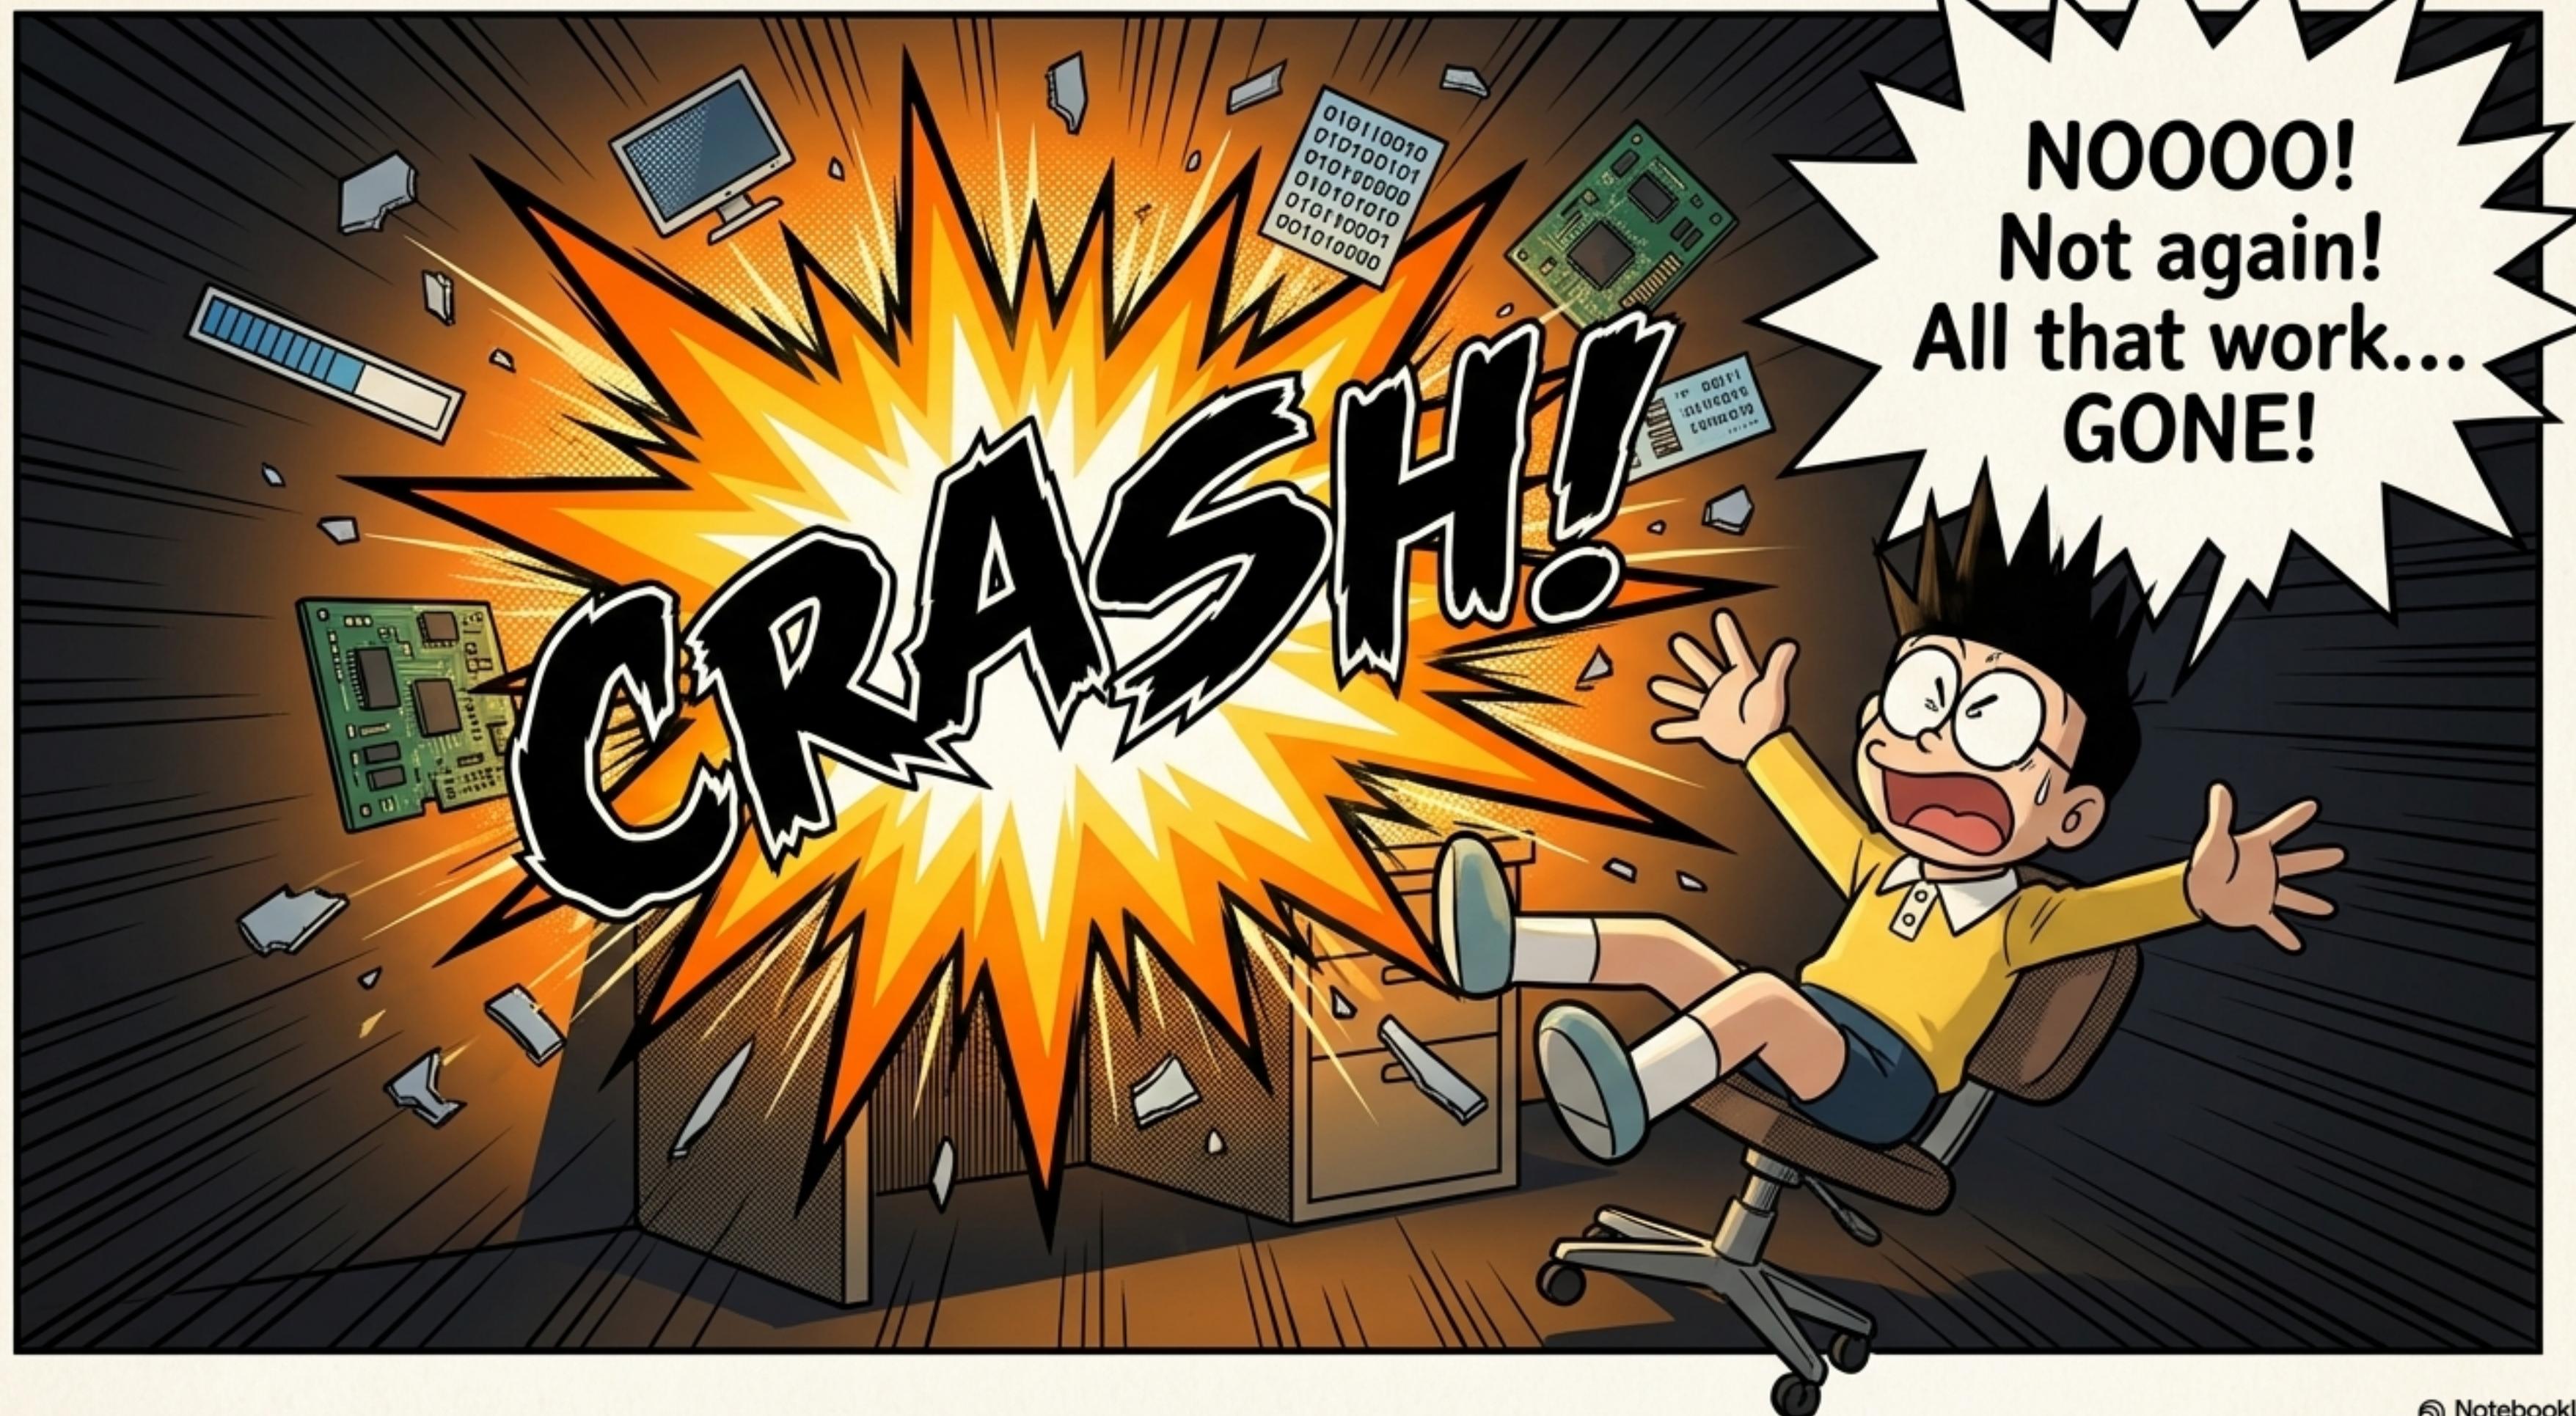


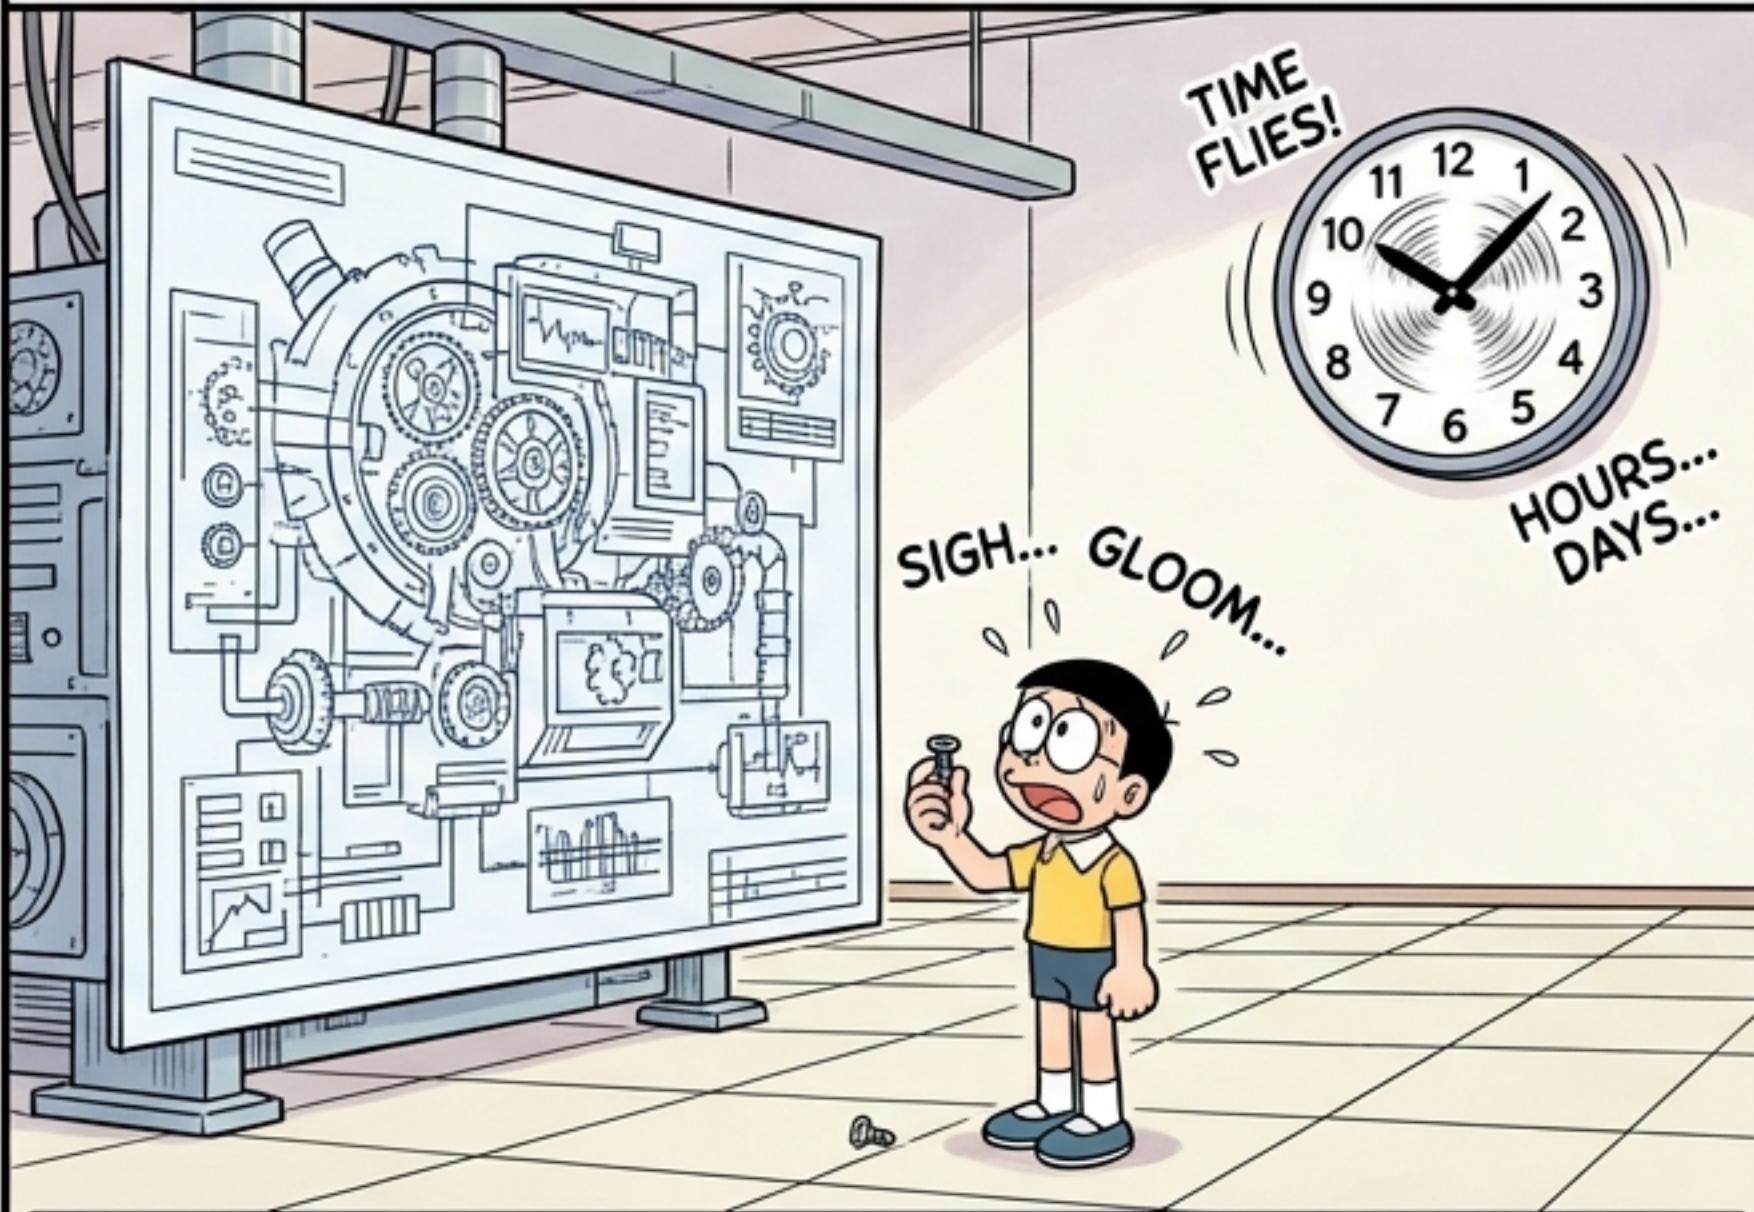
NOOOO!
Not again!
All that work...
GONE!

CRASH!



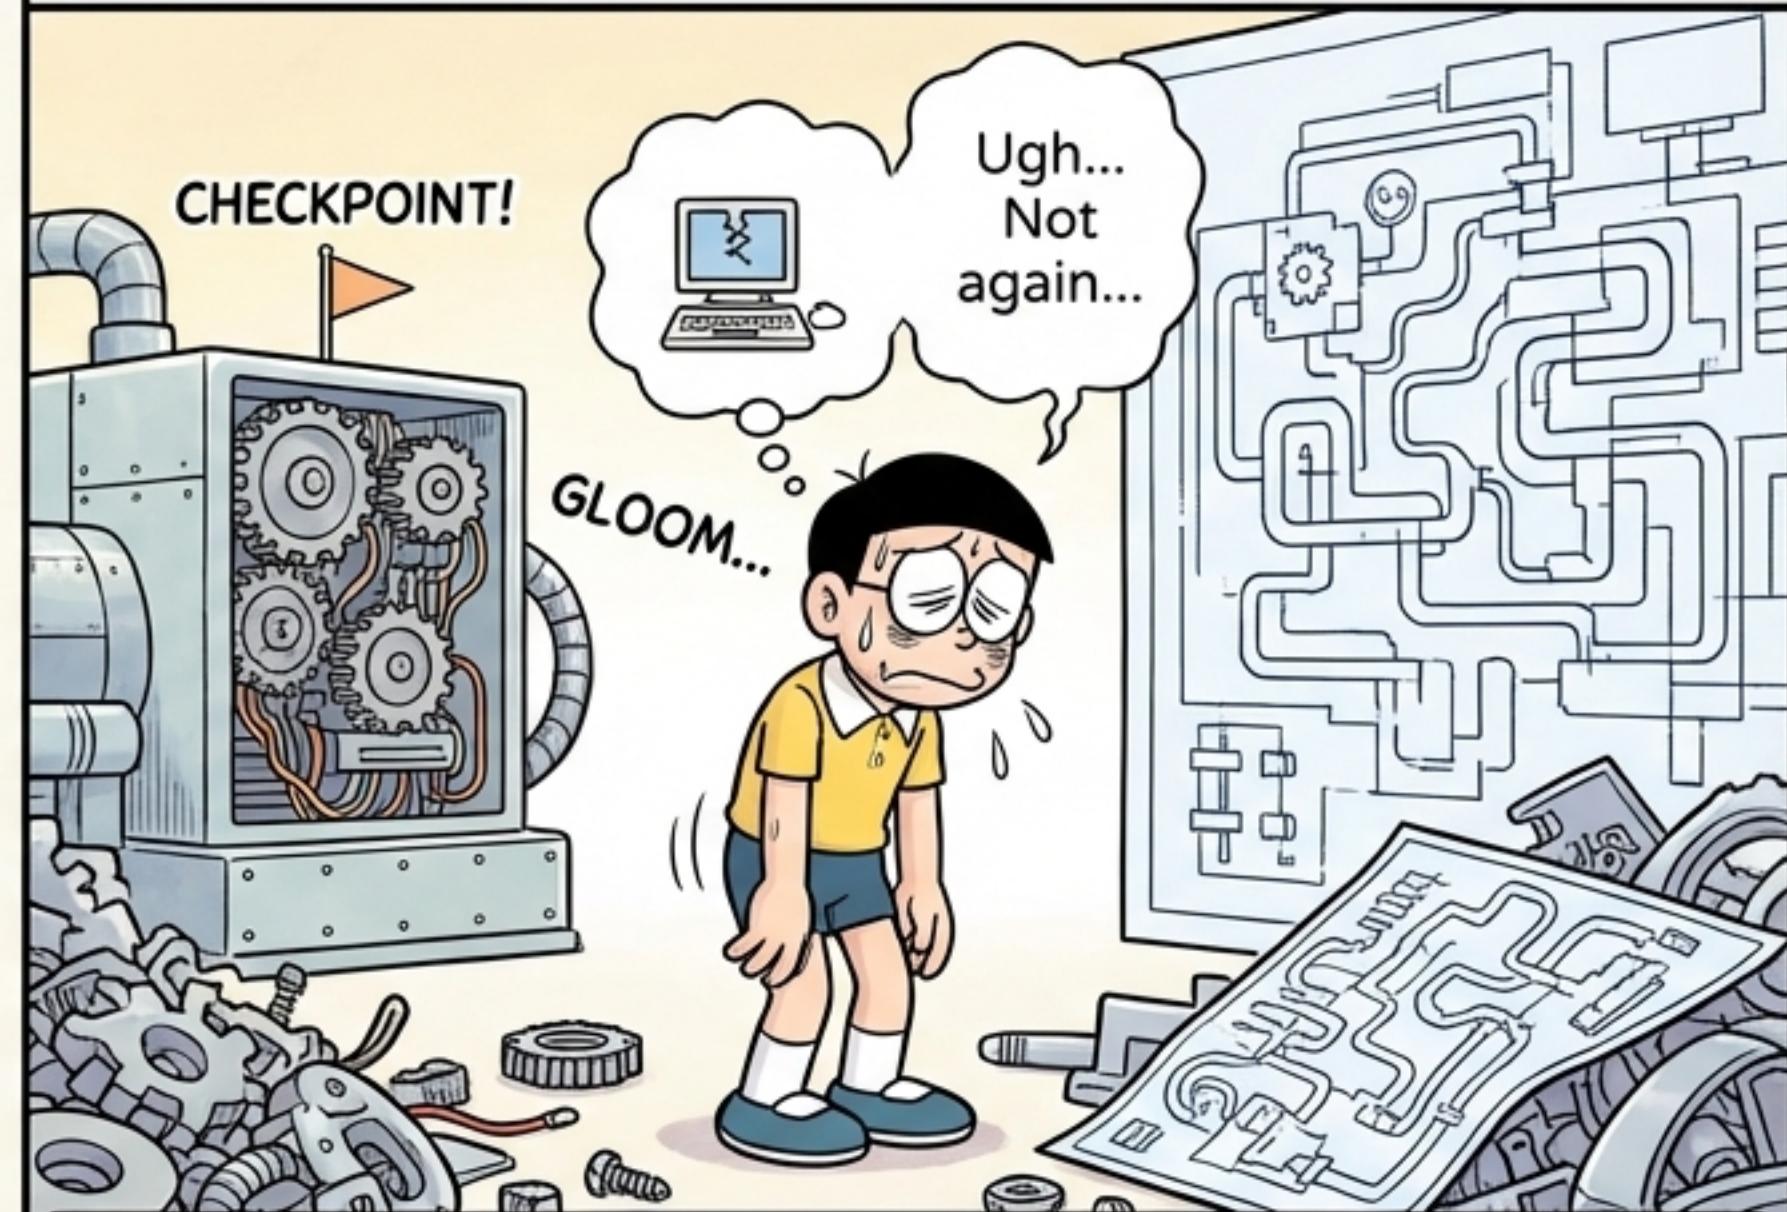
The Old Ways are Painful and Slow.

Paradigm 1: Restart-from-Scratch.

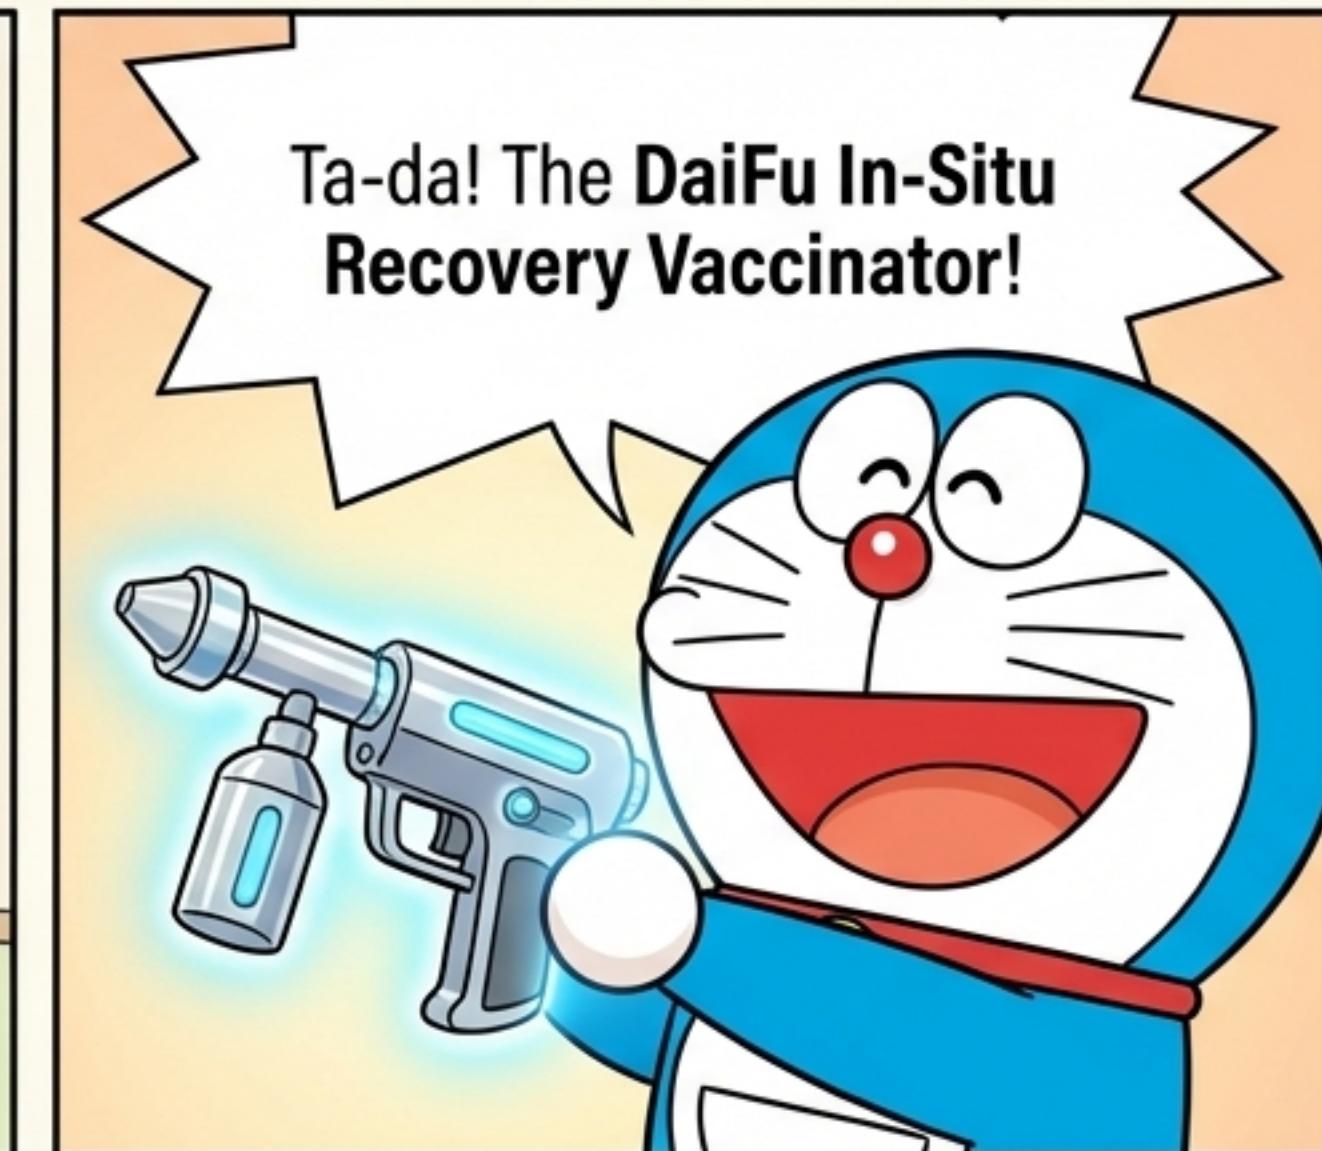
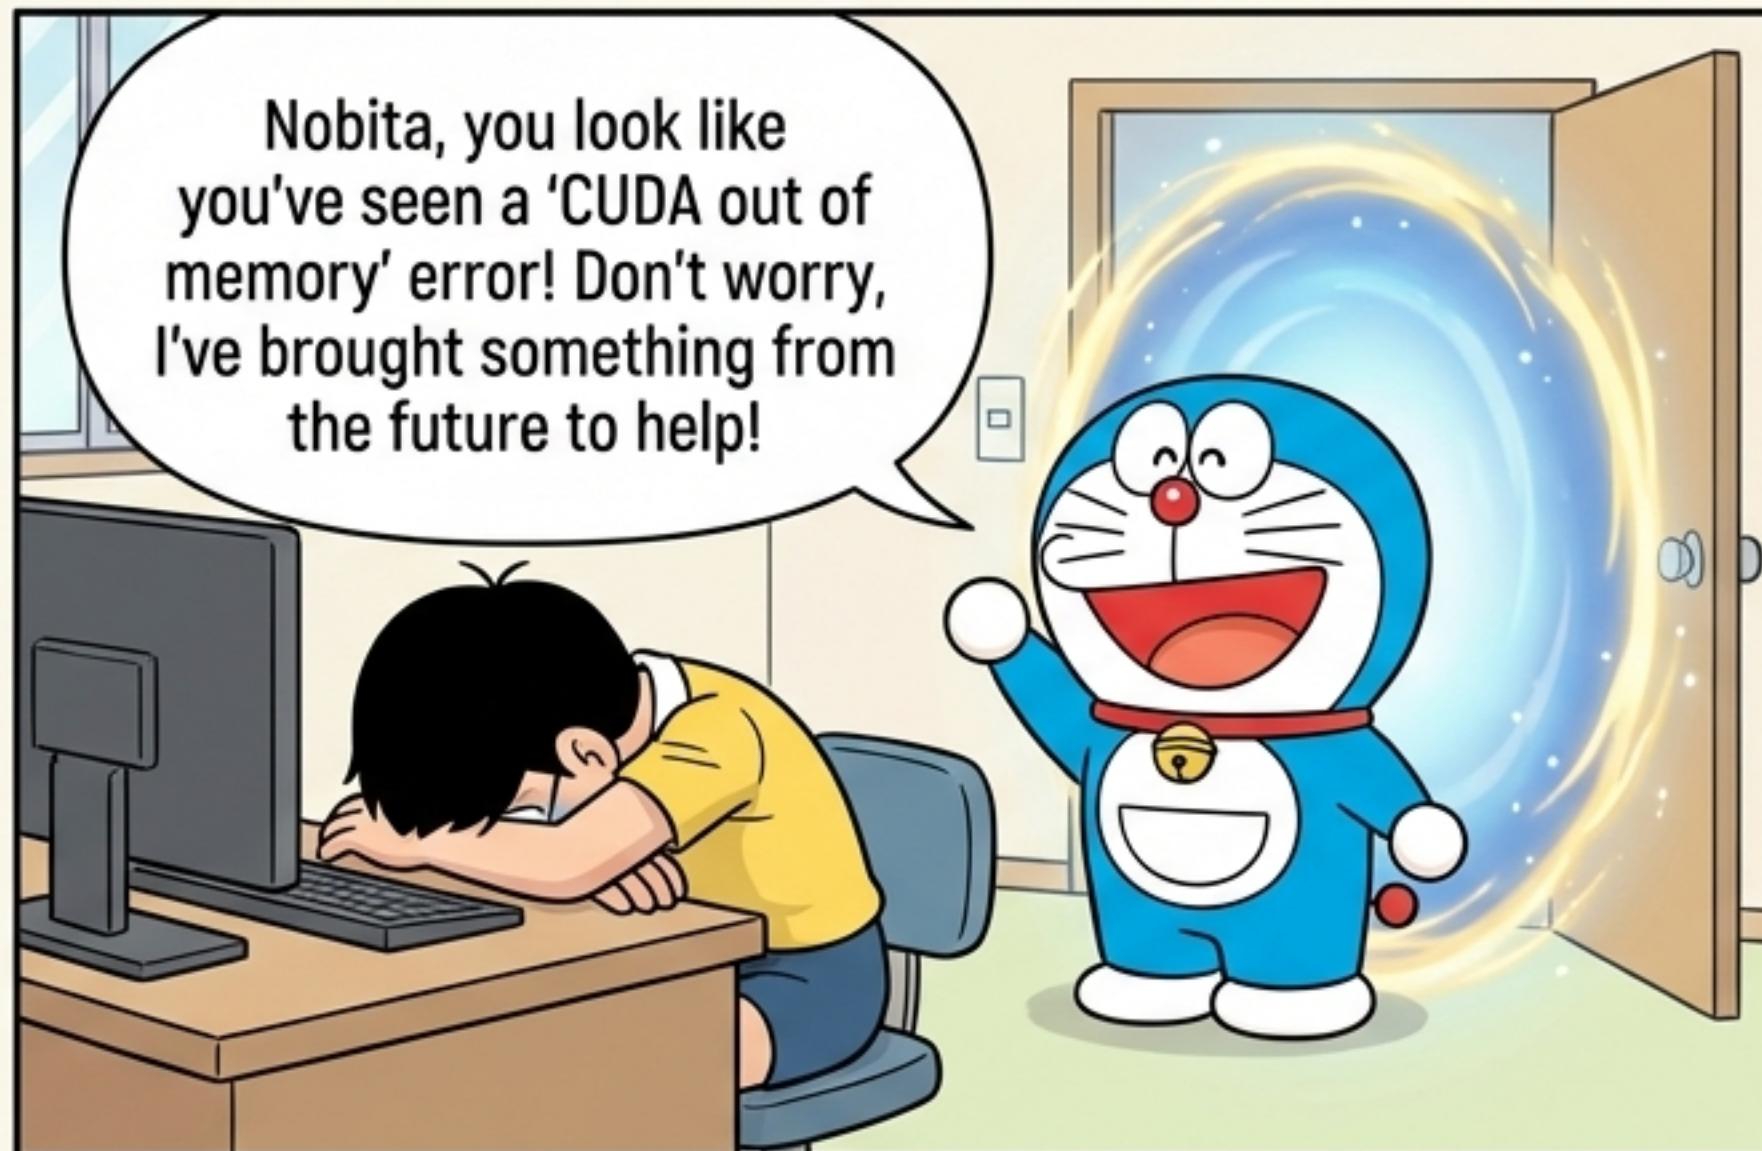


Restarting from the beginning means re-doing hours, or even days, of computation. This wastes huge amounts of resources.

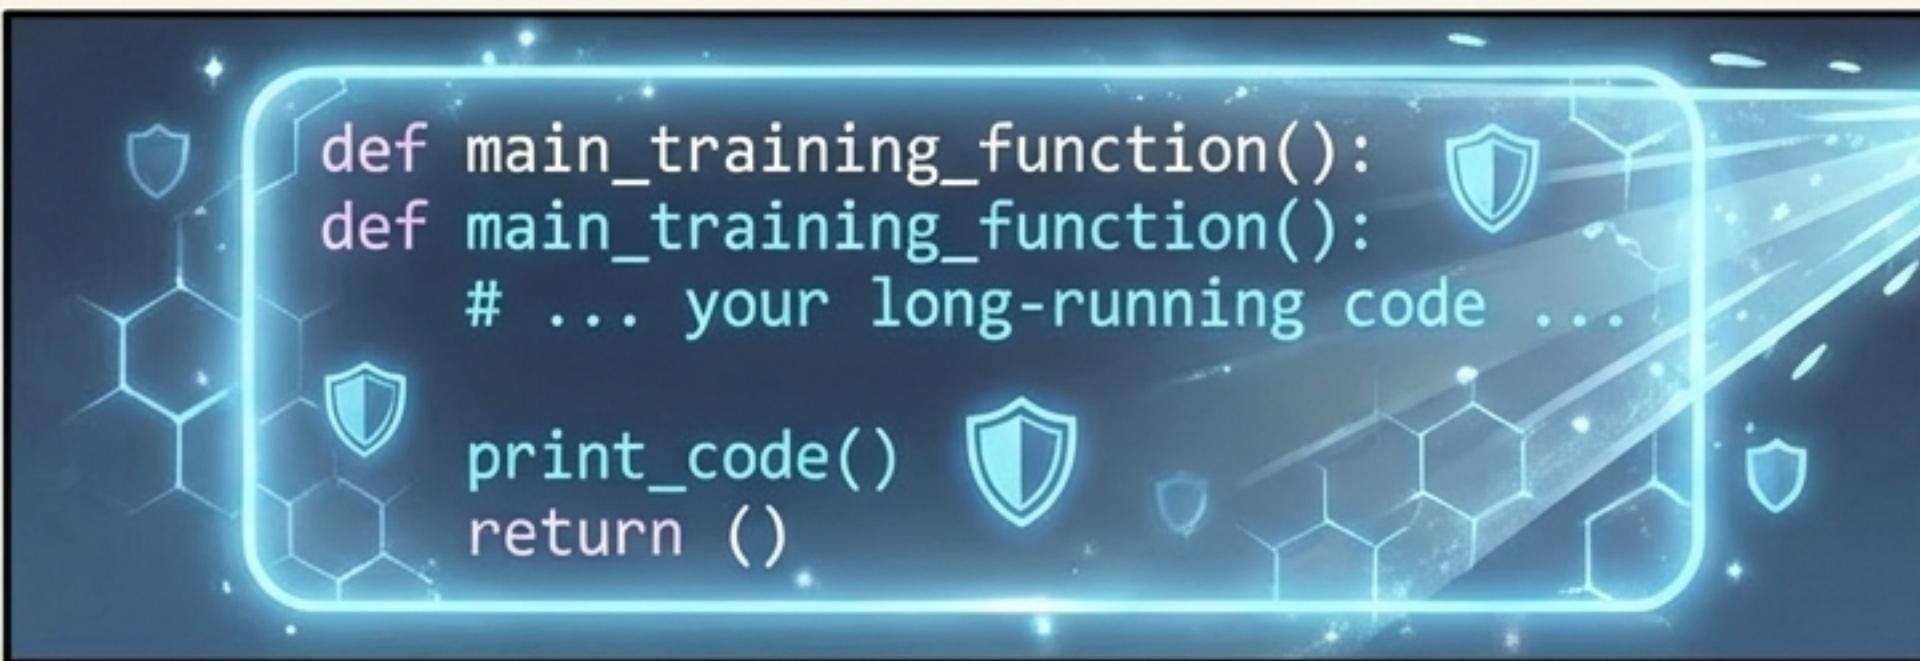
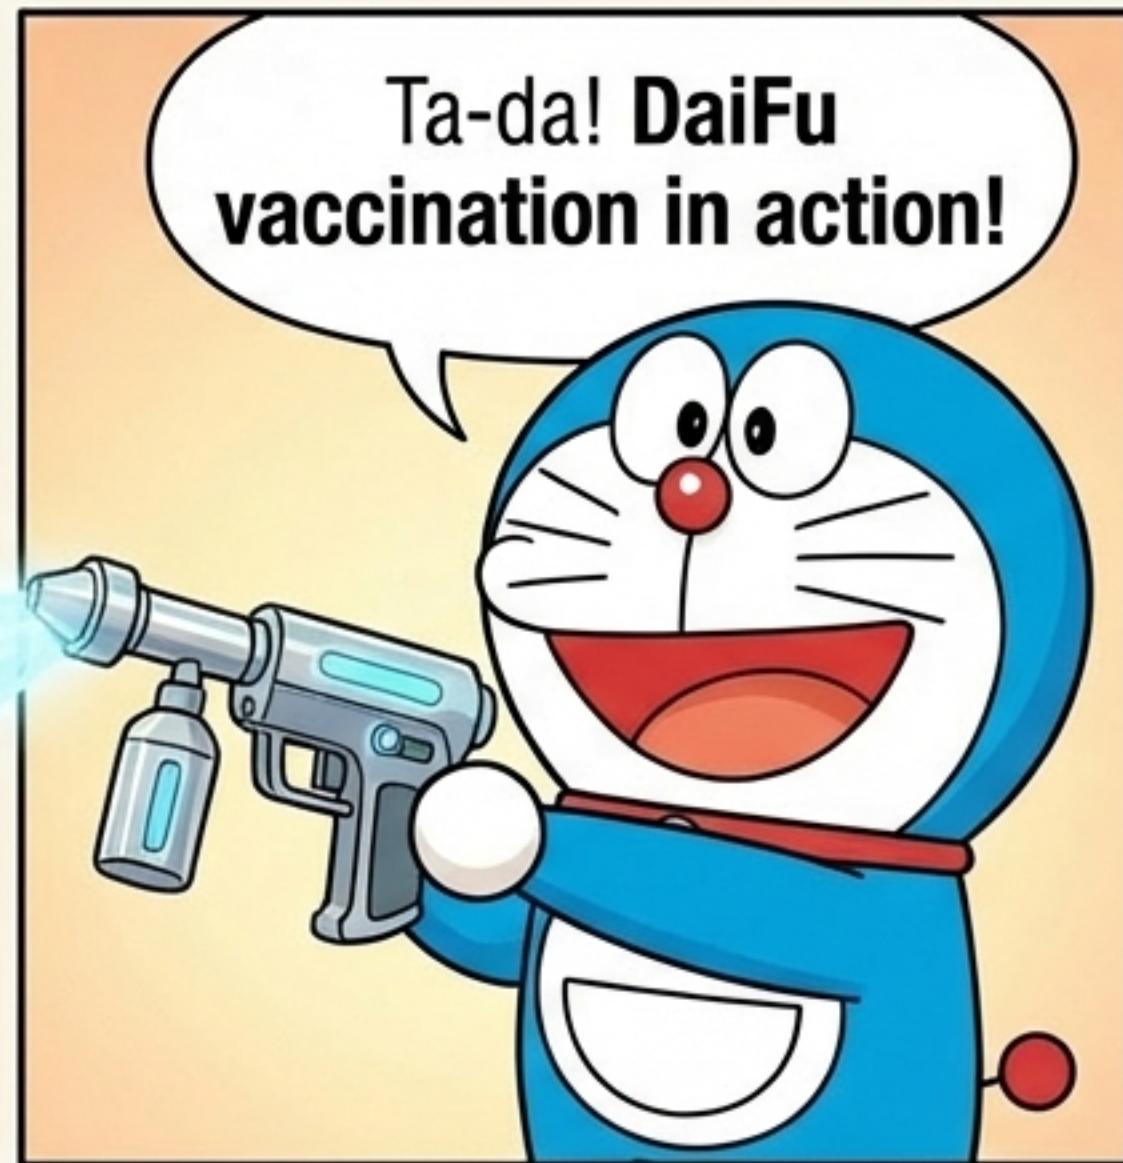
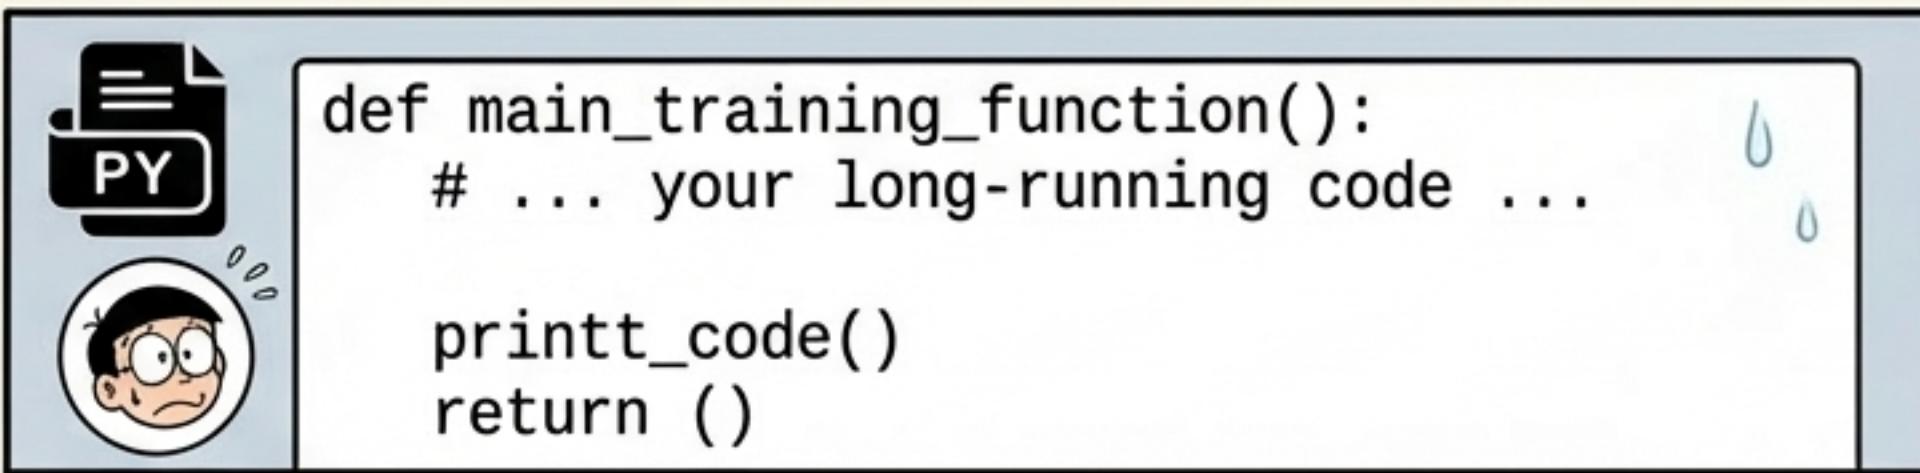
Paradigm 2: Checkpoint-Retry.



Retrying from the last "checkpoint" is better, but you still lose all the progress made since that save point. It's often a heavyweight solution for a minor error.



What is “Program Vaccination”?



DaiFu “vaccinates” a Deep Learning system through a lightweight code transformation. With just two lines of code, it instruments the program to intercept crashes and enable instant, dynamic updates—without needing to restart.

```
import daifu

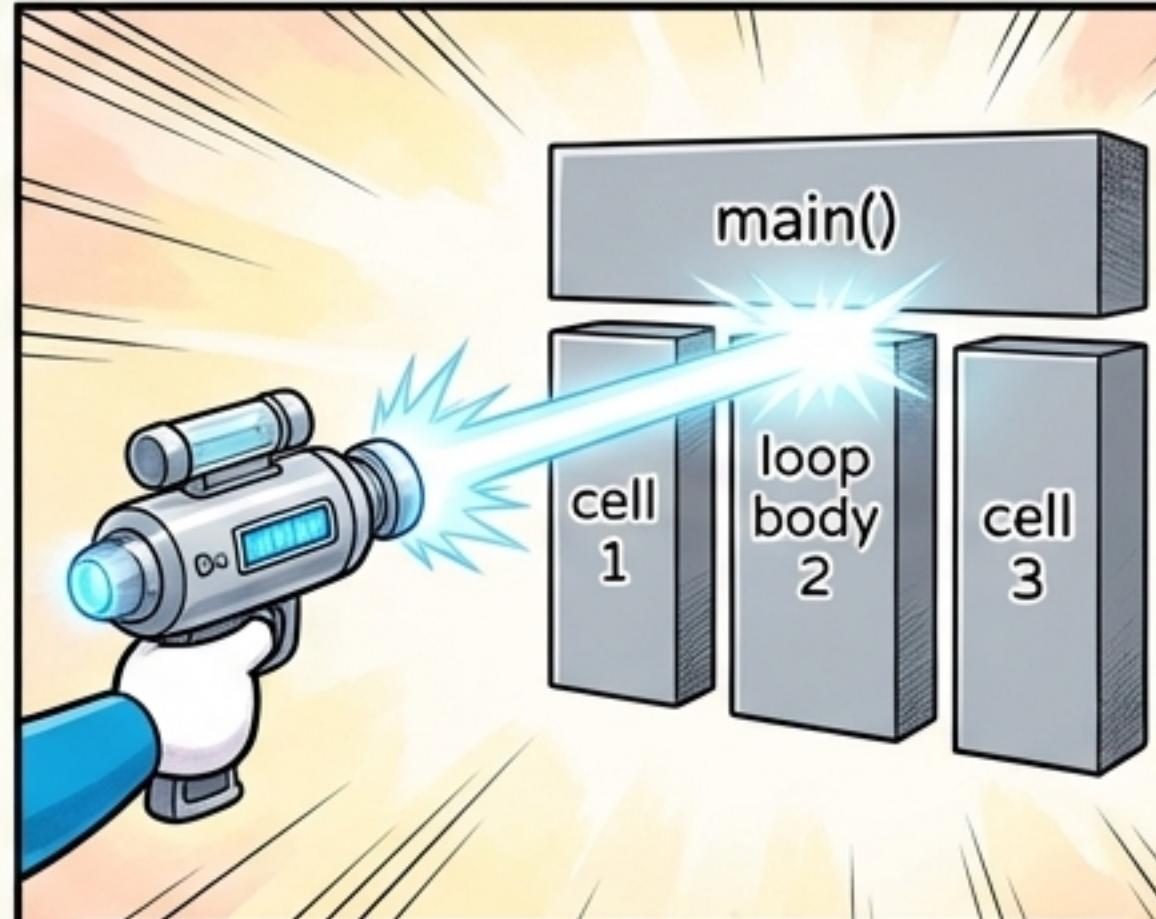
@daifu.transform()
def main_training_function():
    # ... your long-running code ...
```

The Secret: Function Decomposition & Crash Barriers

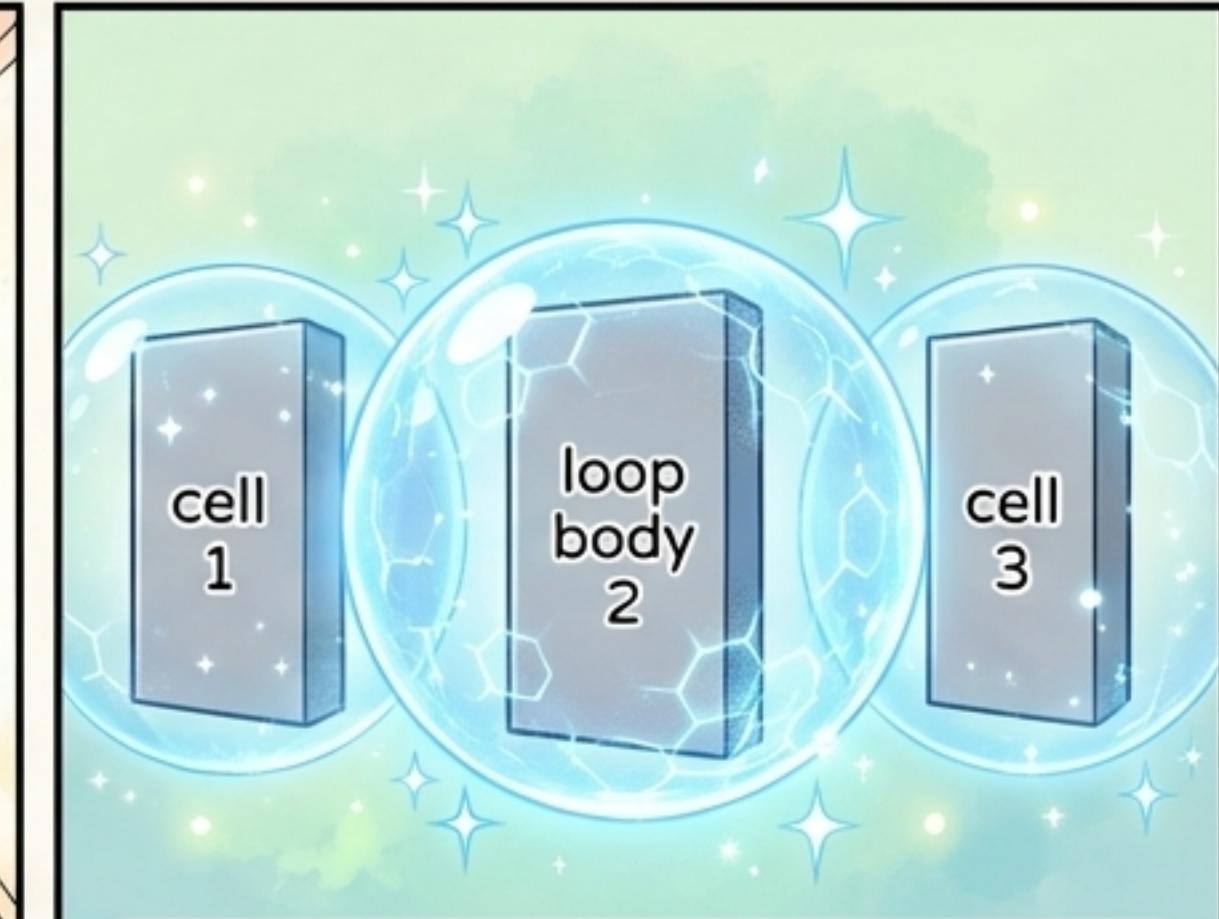
A long-running function is hard to update while it's active.



```
main()
{
    // A large block of abstract code
}
```



****Function Decomposition**:** DaiFu automatically refactors the function into smaller, independent “cells”, often breaking down loops into their own manageable units.

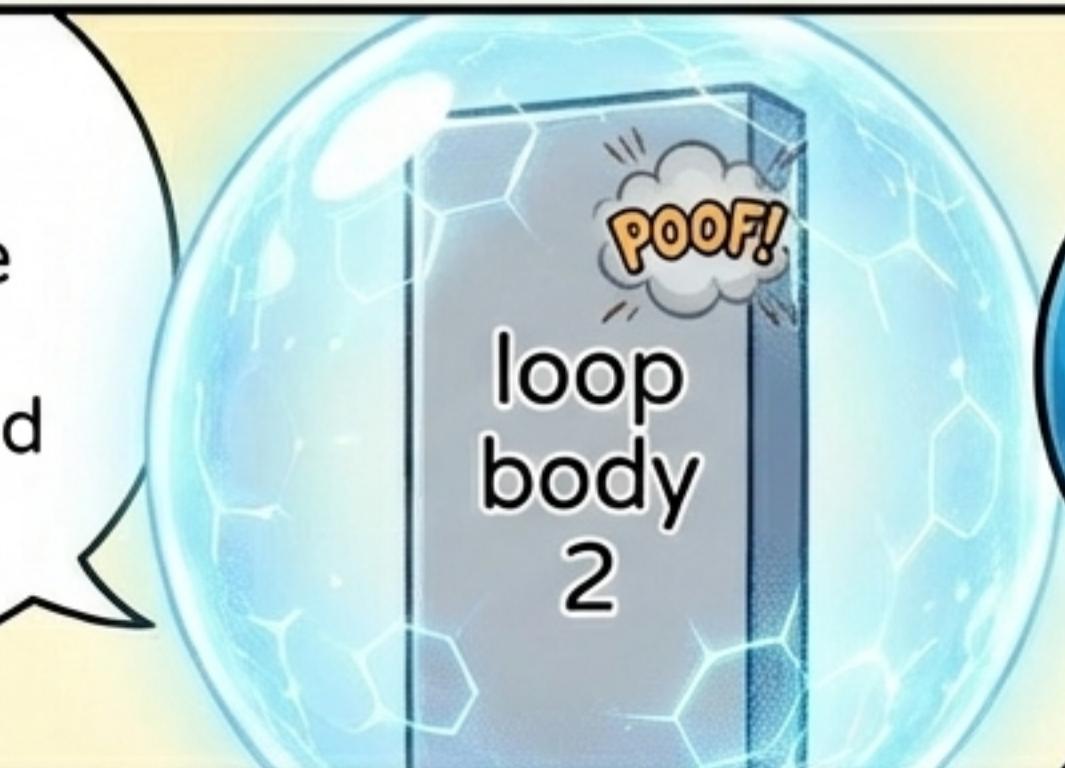


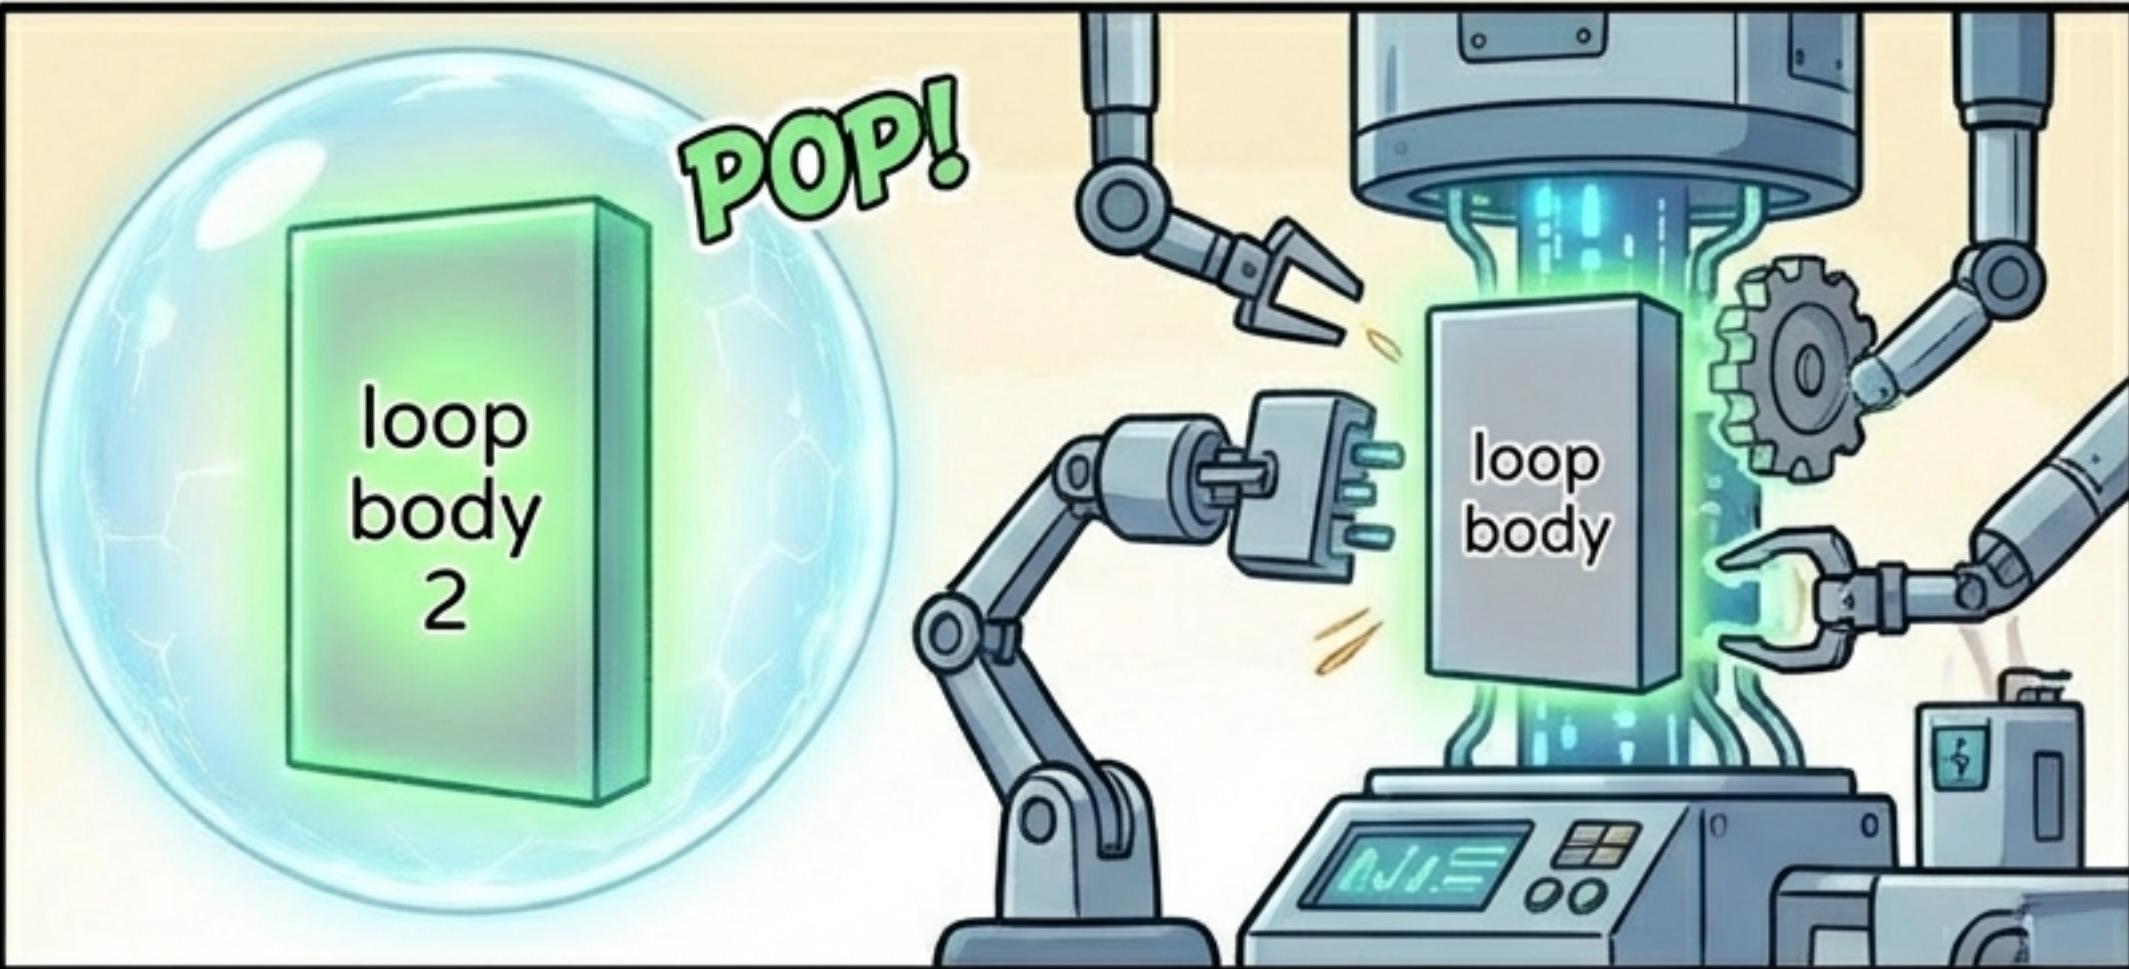
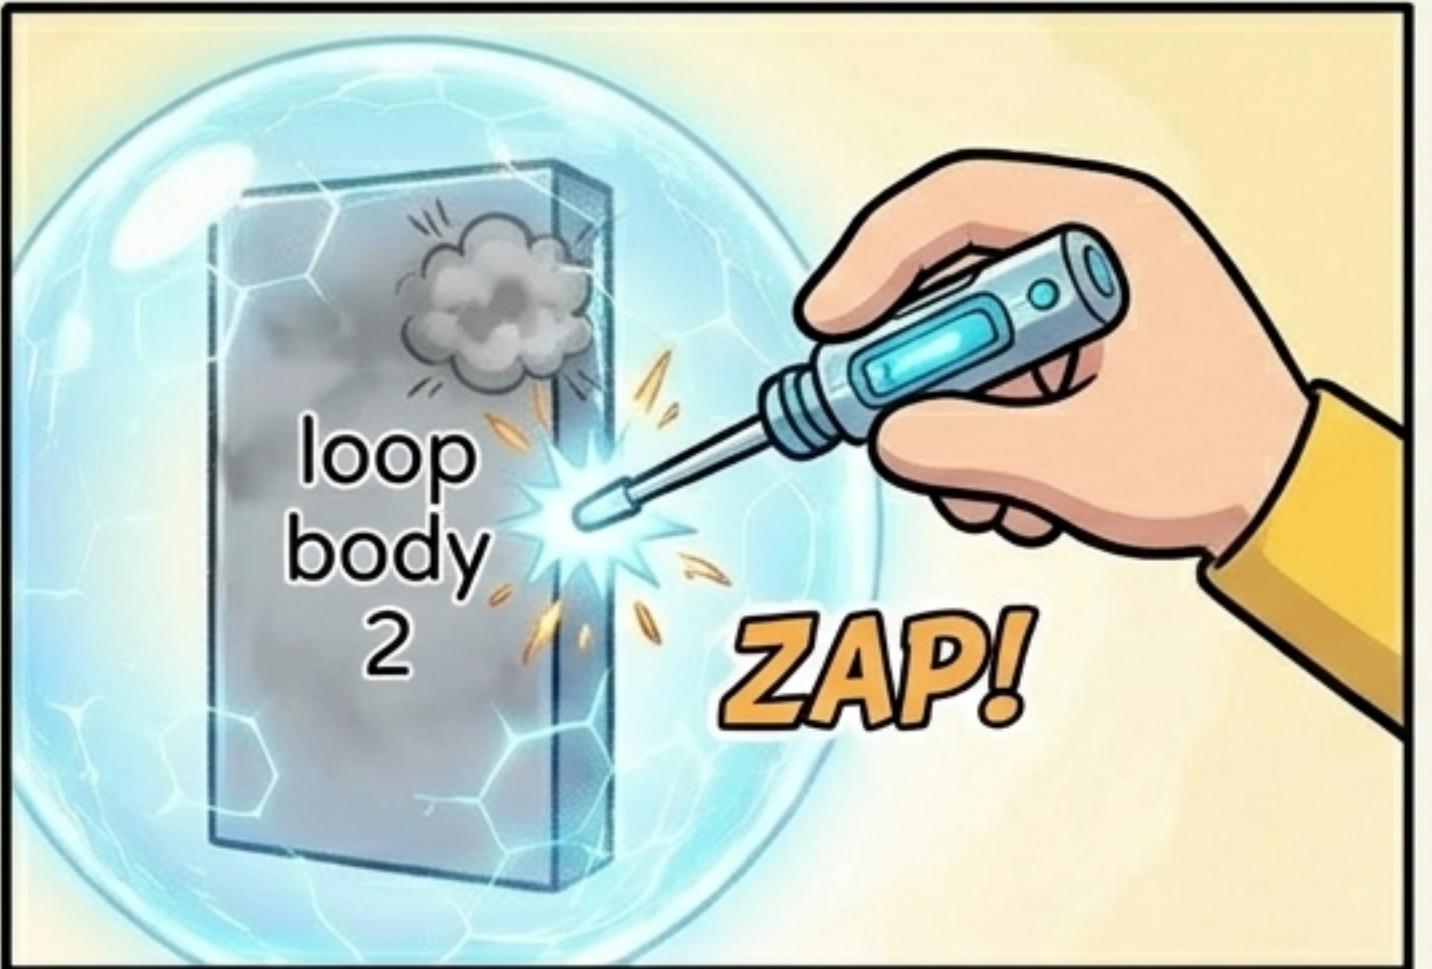
****Crash Barriers**:** Each cell is then wrapped in a “crash barrier”—an exception handler that contains any error within that cell, preventing it from crashing the entire program.

See? The crash barrier contained the error! The rest of your program is safe and still running.

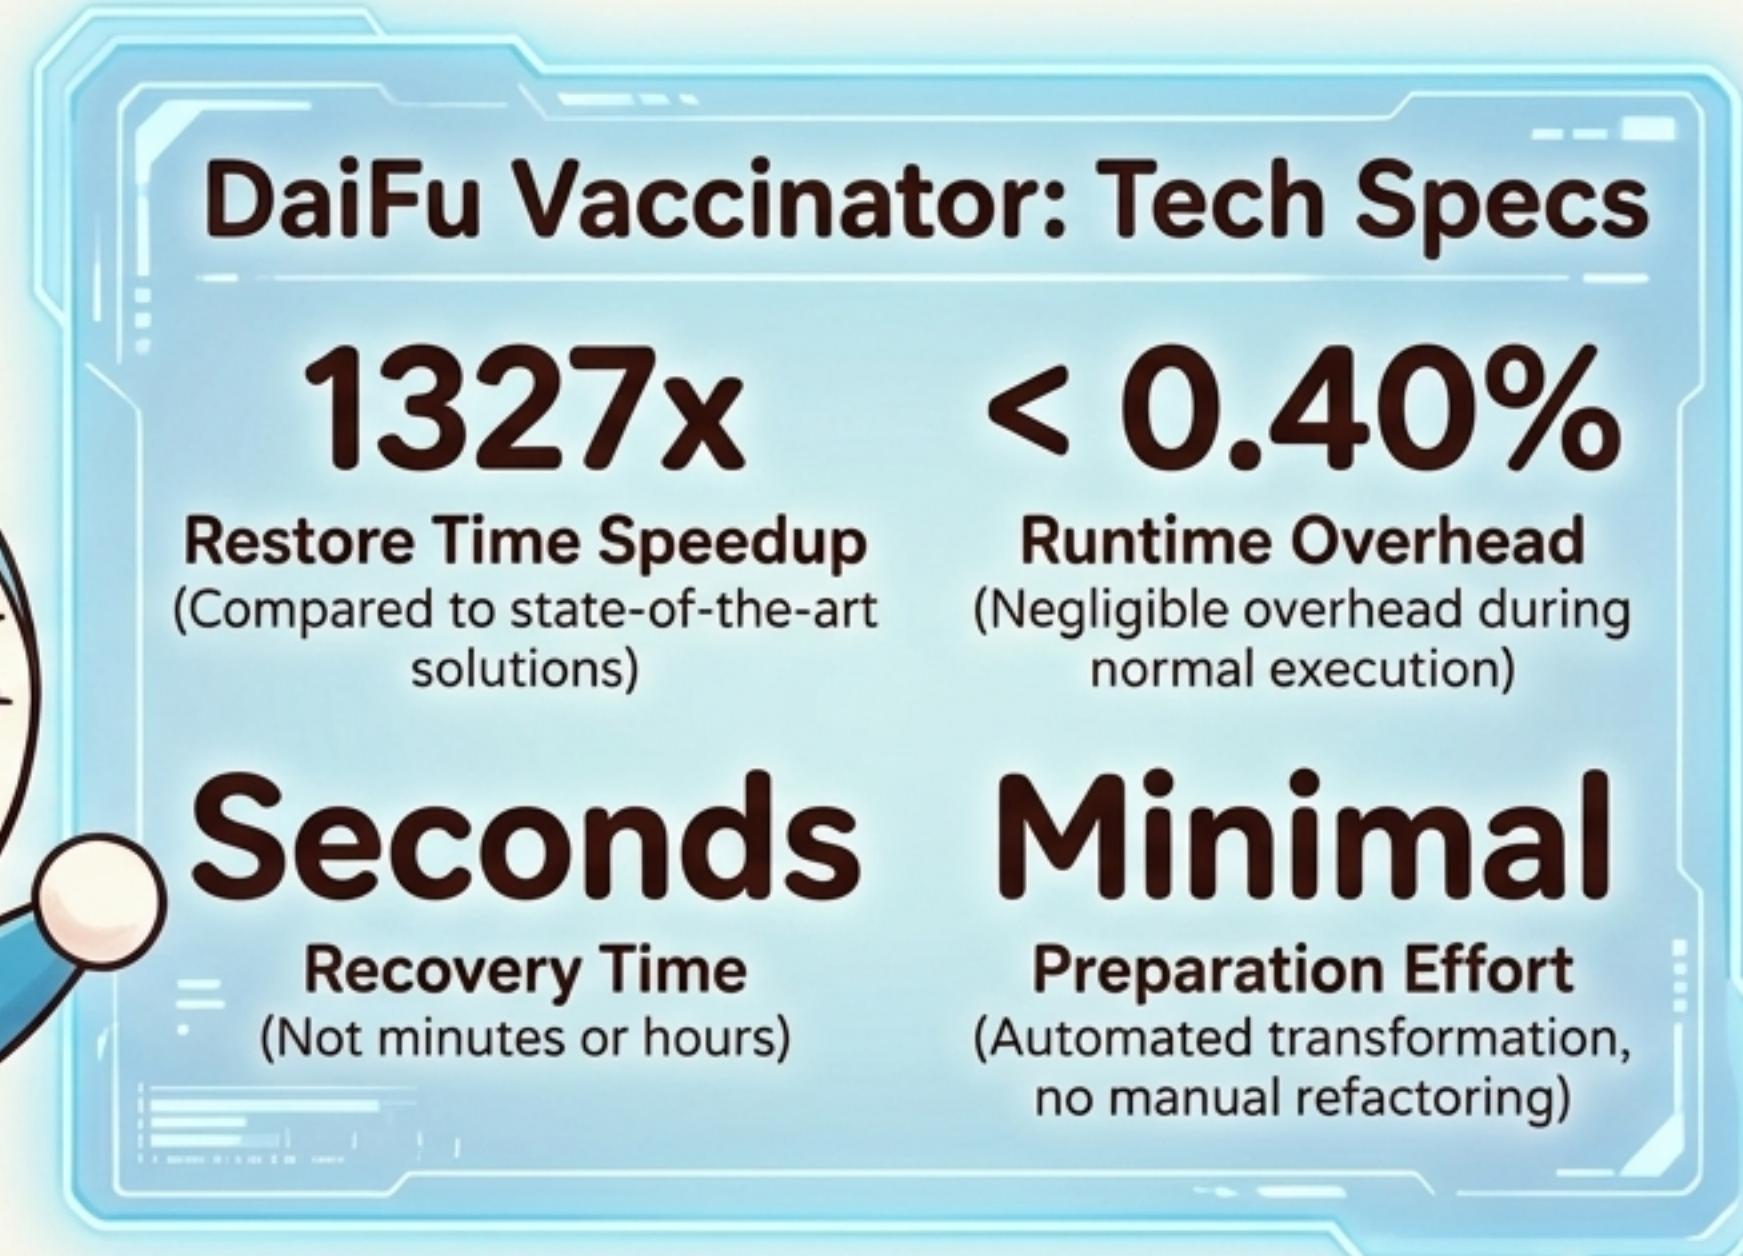


Now for the magic of
In-Situ Recovery!
You can interact with the
program *right now*,
diagnose the problem, and
apply a fix dynamically.





This isn't Magic, It's 22nd Century Engineering!



A Tale of Two Timelines: Crash Recovery Compared

Without DaiFu



Restore Time
— Original — CheckFreq

HOURS



11 12 1 10 2 9 3 8 4 7 6 5 HOURS

With DaiFu

Restore Time
— Original — CheckFreq

30
20
10
0

55 60 50 45 40 35 30 25

SECONDS

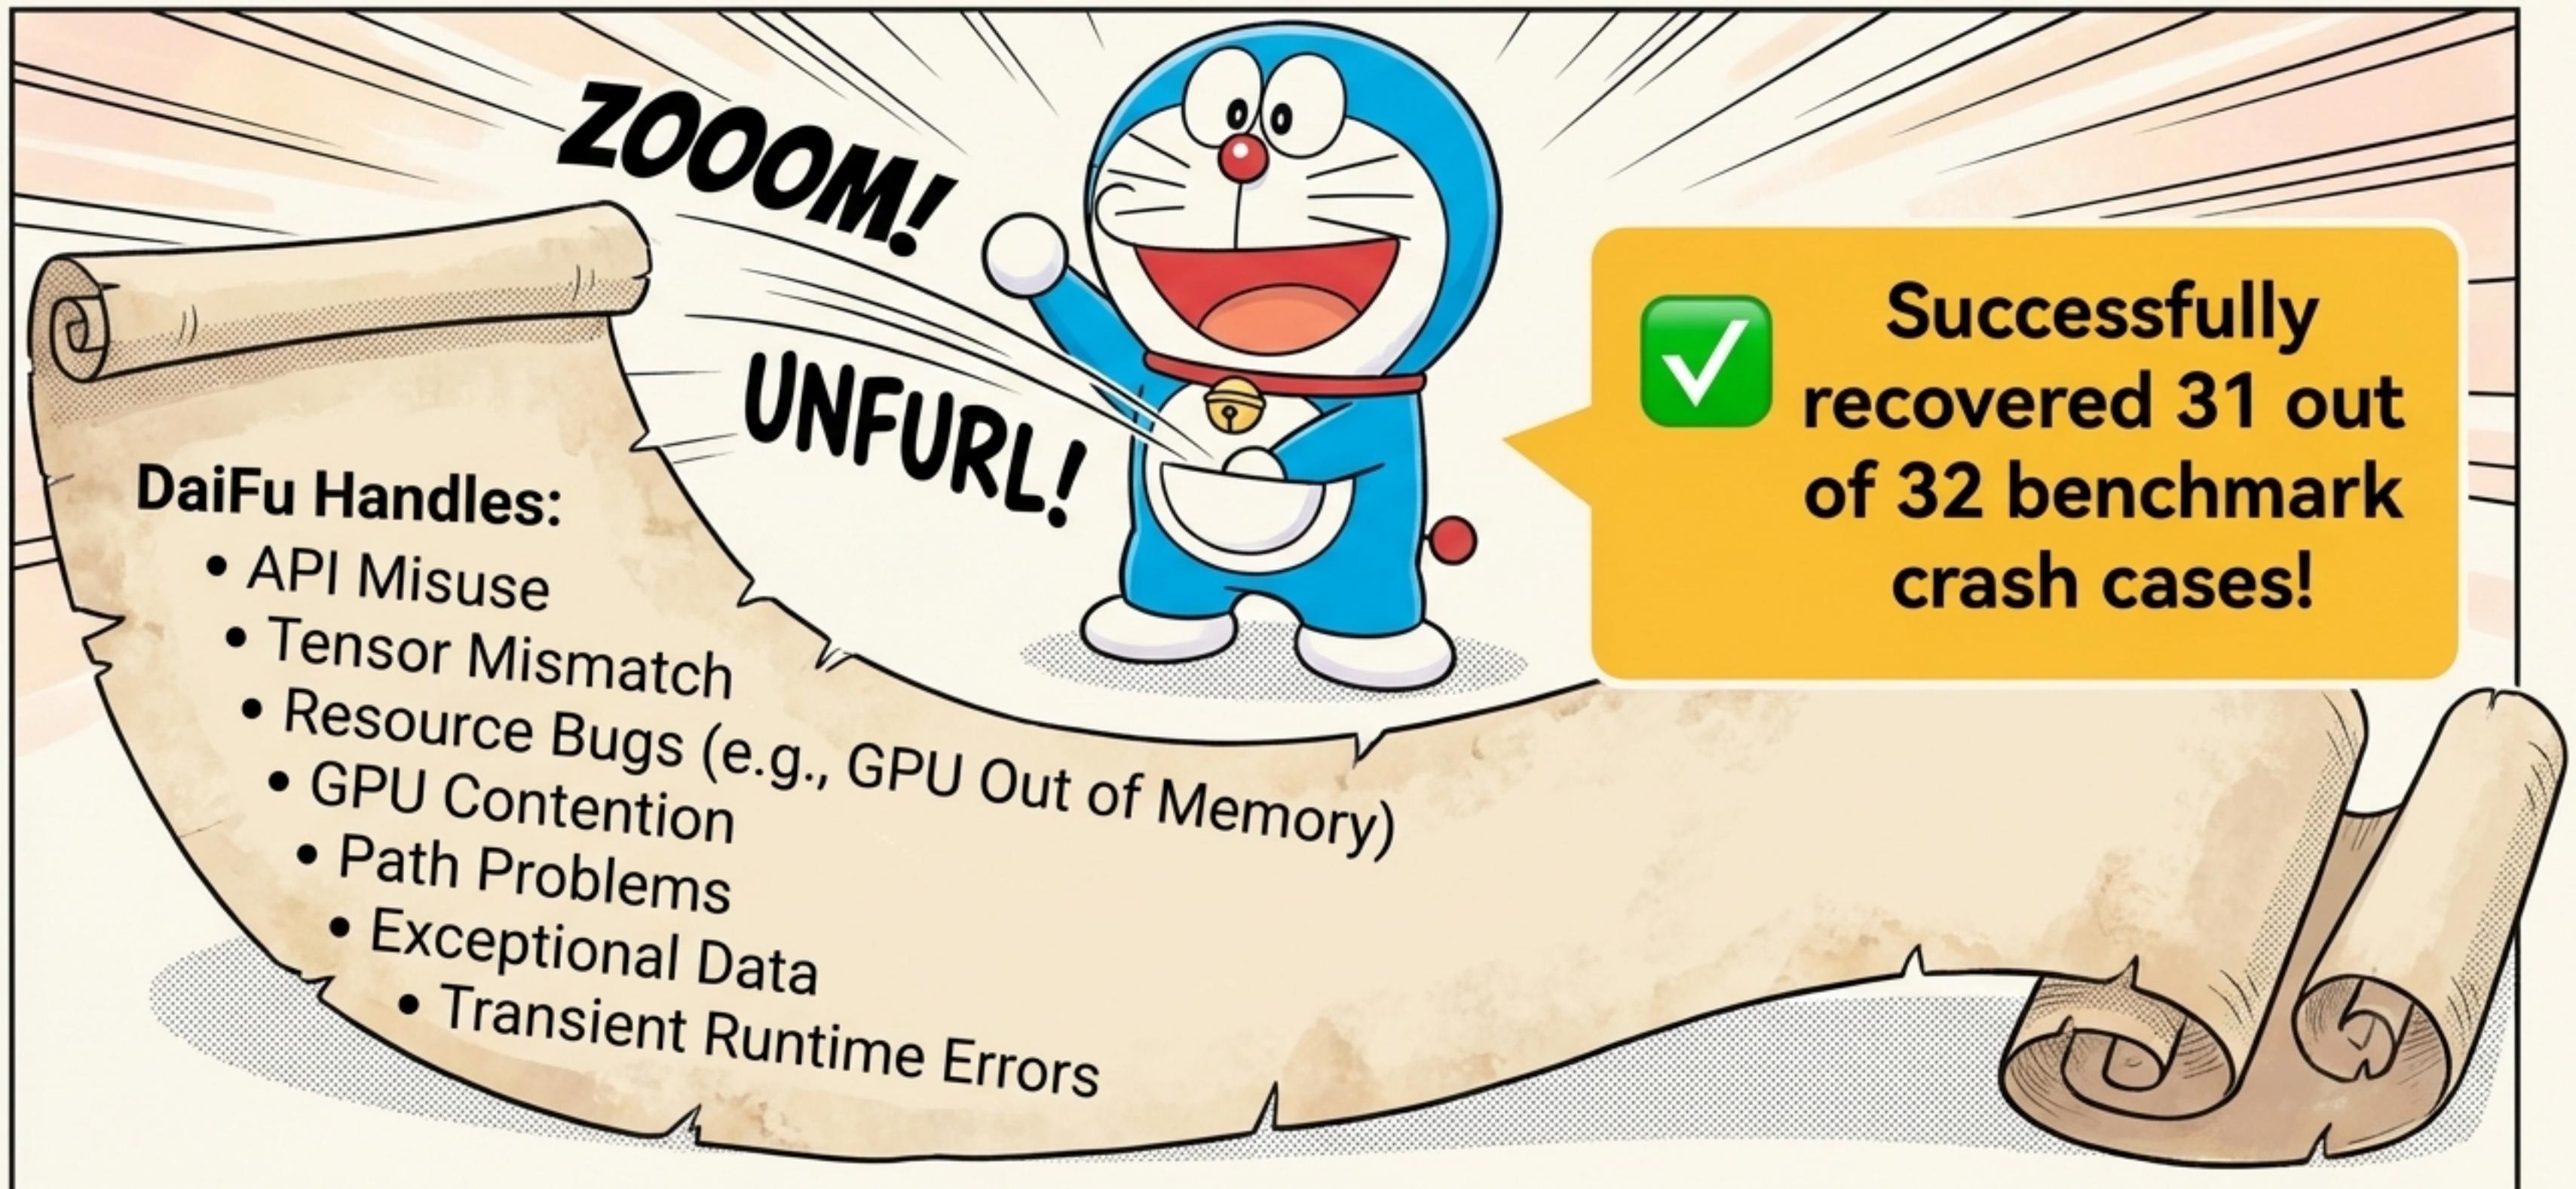
— Original — CheckFreq

— Varela — DaiFu

Wow! It's back already! That was incredibly fast!



Ready for Real-World Problems



DaiFu Handles:

- API Misuse
- Tensor Mismatch
- Resource Bugs (e.g., GPU Out of Memory)
- GPU Contention
- Path Problems
- Exceptional Data
- Transient Runtime Errors



Successfully recovered 31 out of 32 benchmark crash cases!

Case Study: Fixing an “API Misuse” Crash In-Situ

The Bug

```
# The crash occurs here!
test_loss += criterion(output, target).data[0]
```

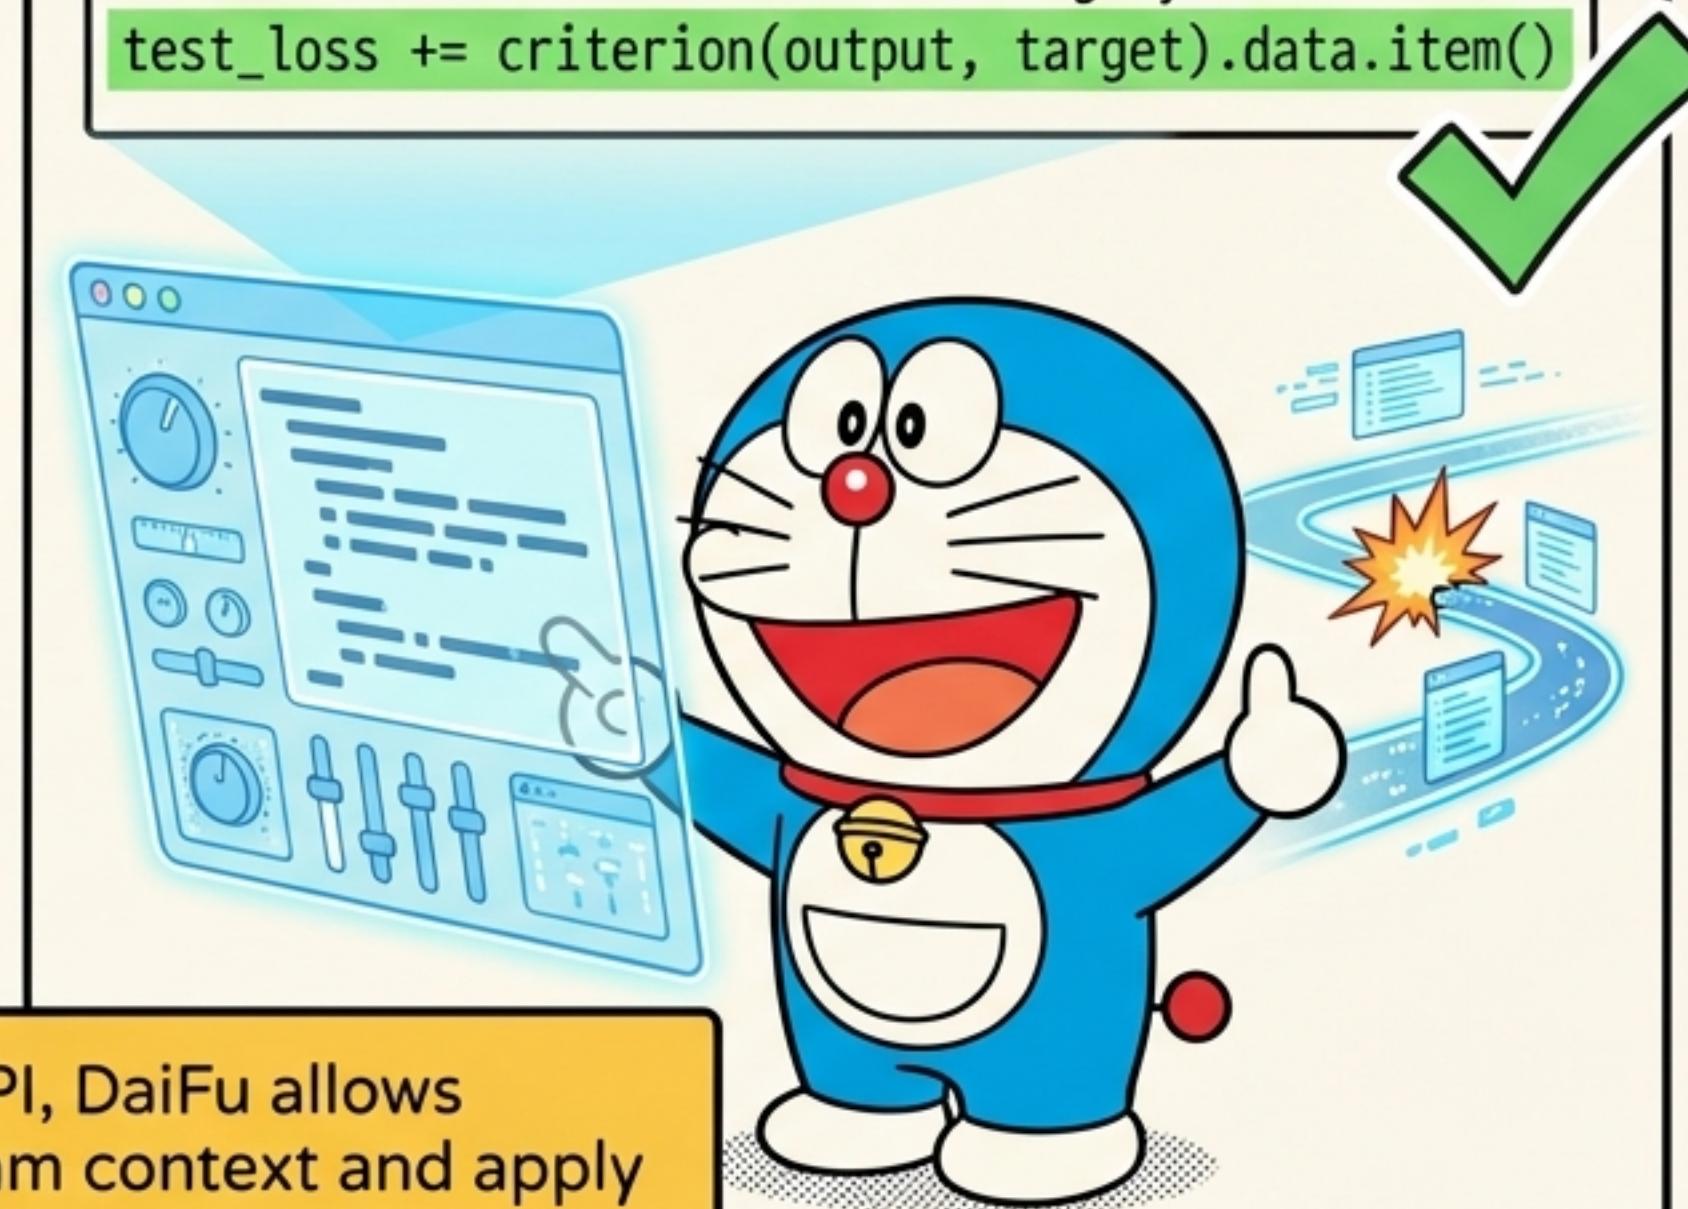
IndexError: invalid index
of a 0-dim tensor.



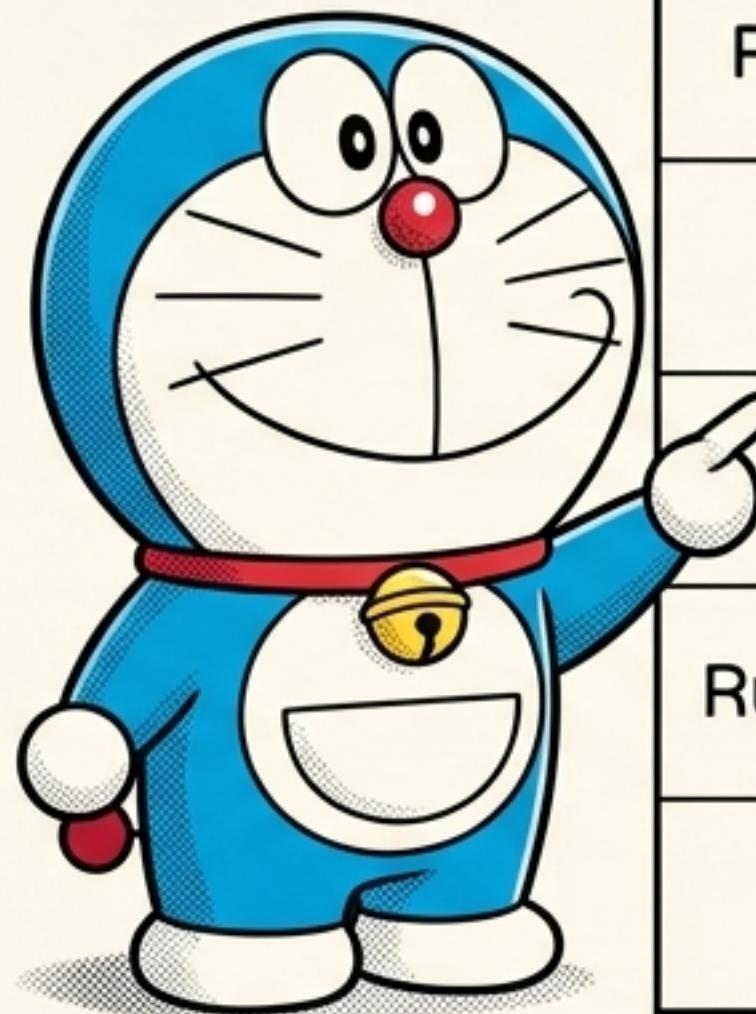
For common programming errors like misusing an API, DaiFu allows developers to diagnose the issue with the live program context and apply a code update instantly, resuming execution without losing state.

The Fix

```
# Fixed in-situ with DaiFu's 'surgery' interface!
test_loss += criterion(output, target).data.item()
```



DaiFu: A New Paradigm for Crash Recovery



Feature	Restart / Checkpoint-Retry	DaiFu (In-Situ Recovery)
Retry Granularity	Program / Epoch / Iteration	Specific Lines of Code
Code Updates	Requires Full Program Restart	Dynamic, Instant Updates
Restore Time	Minutes to Hours	Seconds
Runtime Overhead	Can be high (e.g., 7.64%)	Negligible (<0.40%)
User Effort	Manual code refactoring	Automated (2-line decorator)

DaiFu is not just an optimization of checkpointing; it's a fundamentally different and more agile approach to handling failures in DL systems.

The Science Behind the 22nd Century Gadget



Perspective

Introduces and improves Dynamic Software Updating (DSU) to recover crashes in DL systems *in situ*—a new paradigm that achieves almost instant restoration.



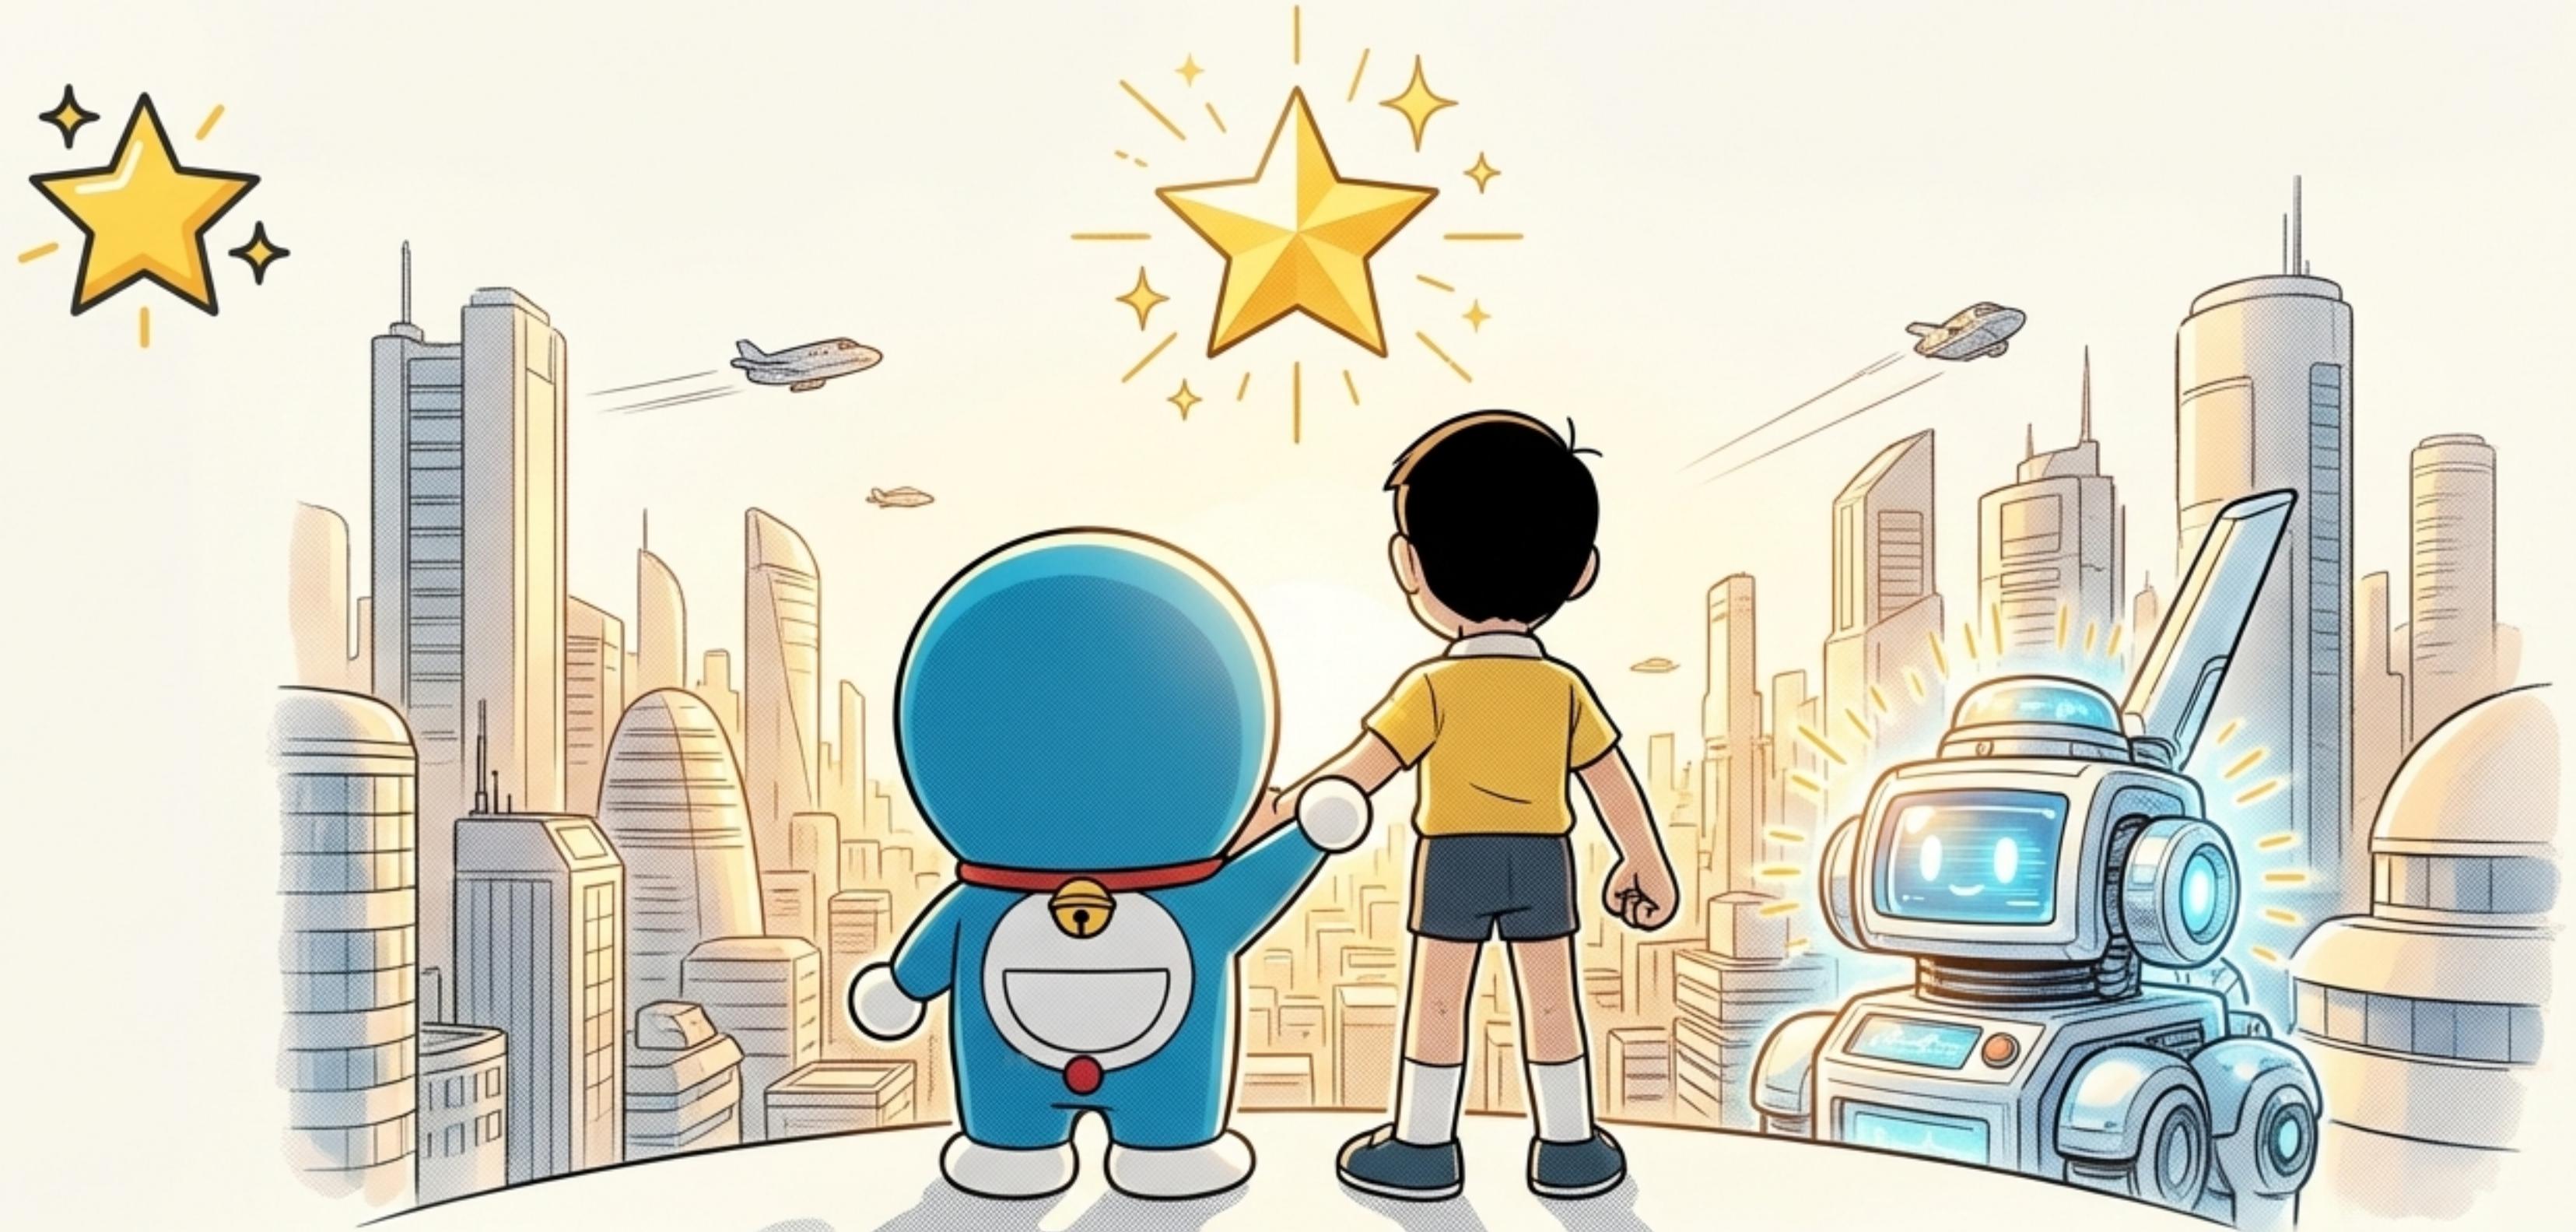
Framework

Proposes the novel DaiFu framework, which uses “Program Vaccination” to enable quick recovery while incurring low runtime overhead and minimal preparation effort.



Benchmark

Contributes a reproducible benchmark of 32 crash cases across 7 distinct scenarios, enabling comprehensive evaluation of DL crash recovery techniques.



By treating crashes not as catastrophic failures but as interceptable events, DaiFu turns hours of frustrating downtime into seconds of productive problem-solving. This accelerates development, saves vast computing resources, and empowers the next generation of AI innovation.