

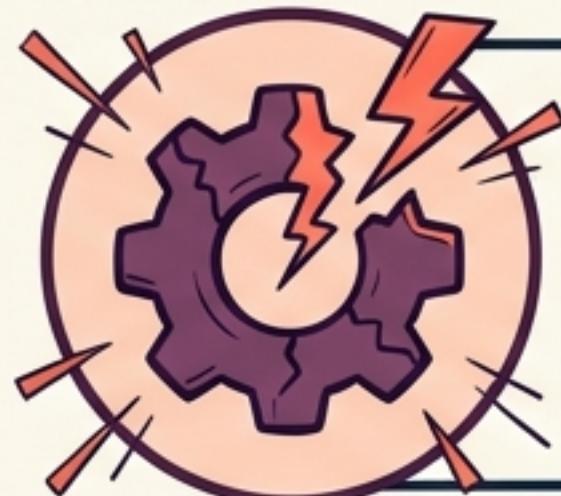
# THE MYSTERY OF THE FLAKY CI JOBS

It passed just  
a minute ago!  
Why did it fail  
it \*now\*?!



A TELUS Adventure in Diagnosing and Defeating  
Non-Deterministic Pipeline Failures

# OUR MISSION BRIEFING



## 1. THE CI/CD BATTLEFIELD

A developer's daily struggle with mysterious pipeline failures.



## 2. UNMASKING THE ENEMY

Introducing a systematic approach to identify and categorize every type of flaky failure.



## 3. CALCULATING THE DAMAGE

Quantifying the real cost of these failures in time and resources.



## 4. TRACKING THEIR MOVEMENTS

Analyzing failure patterns to see which are old news and which are new threats.



## 5. THE ULTIMATE GADGET

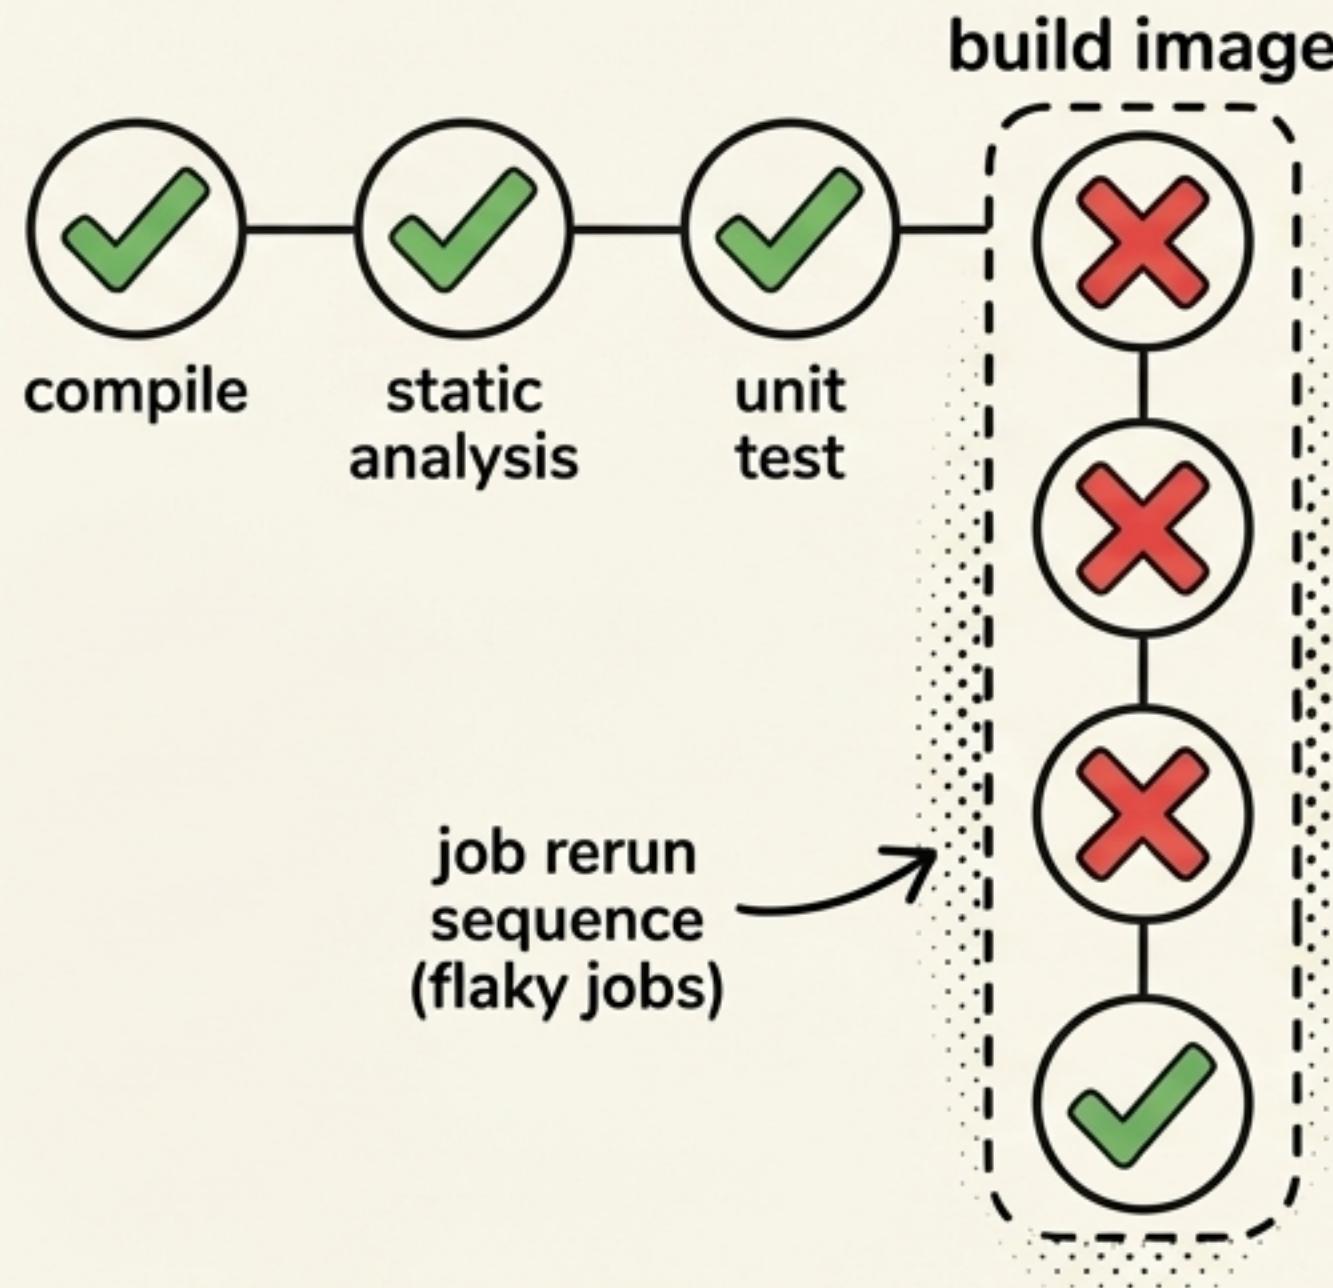
Unveiling a powerful model to prioritize the most critical failures.



## 6. THE BATTLE PLAN

Our final, data-driven strategy to conquer flaky failures.

# A Quarter of Our CI Failures are Ghosts in the Machine



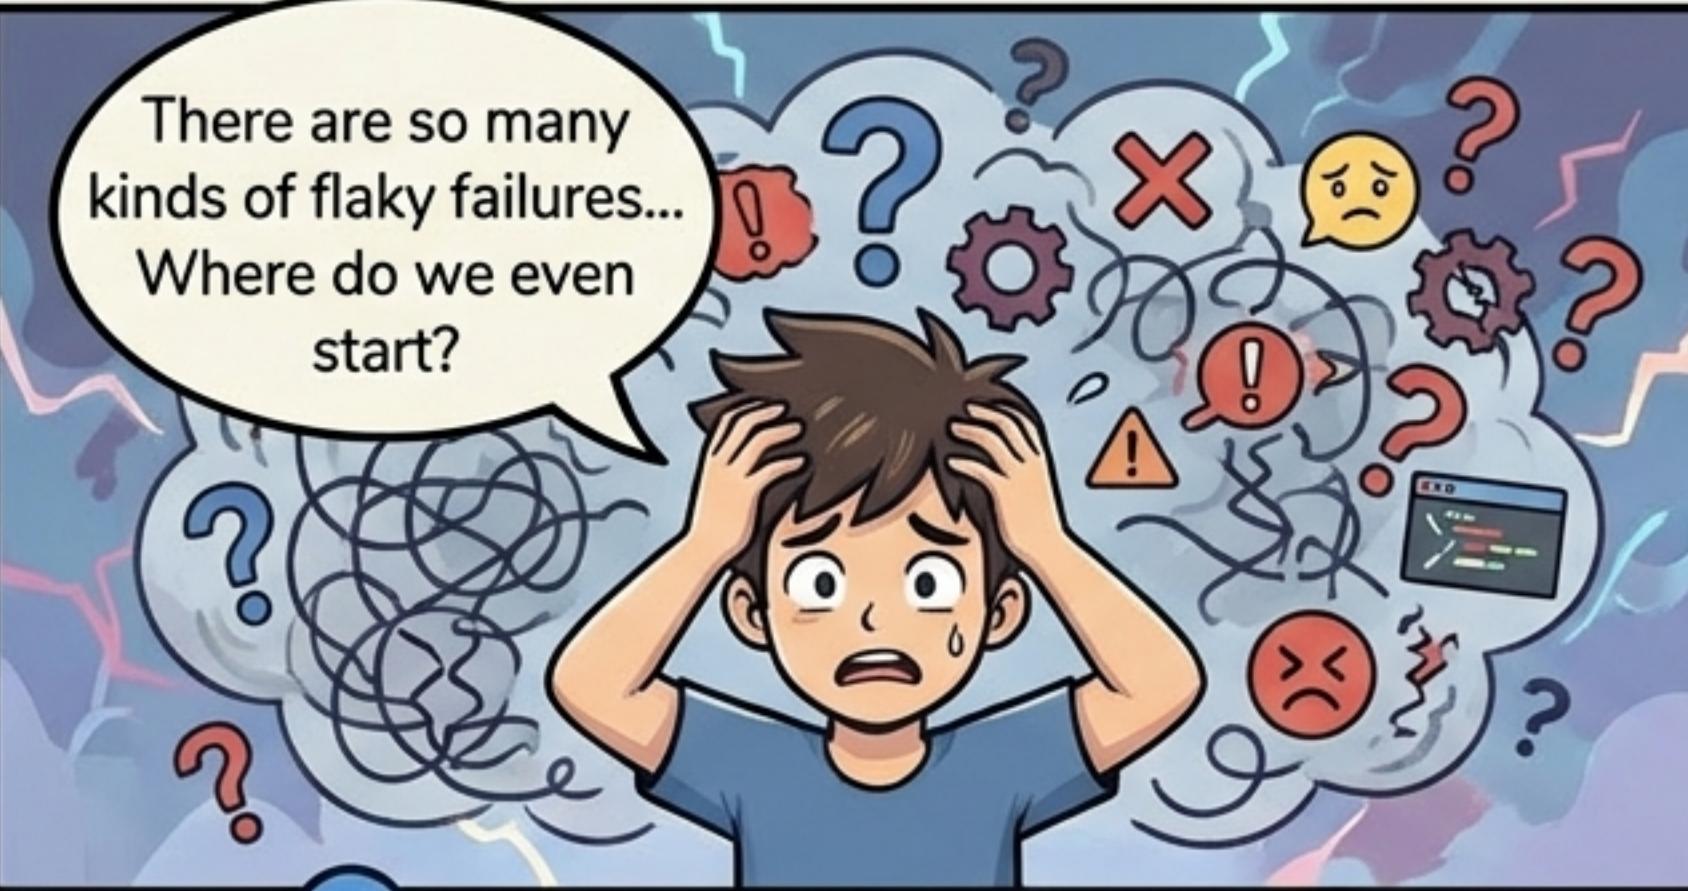
At TELUS, Continuous Integration (CI) pipelines are crucial for rapid software delivery. But they are haunted by “flaky” failures—jobs that fail unpredictably for reasons unrelated to code changes.

**1 in 4  
(25%)**  
of all failed jobs at TELUS is flaky

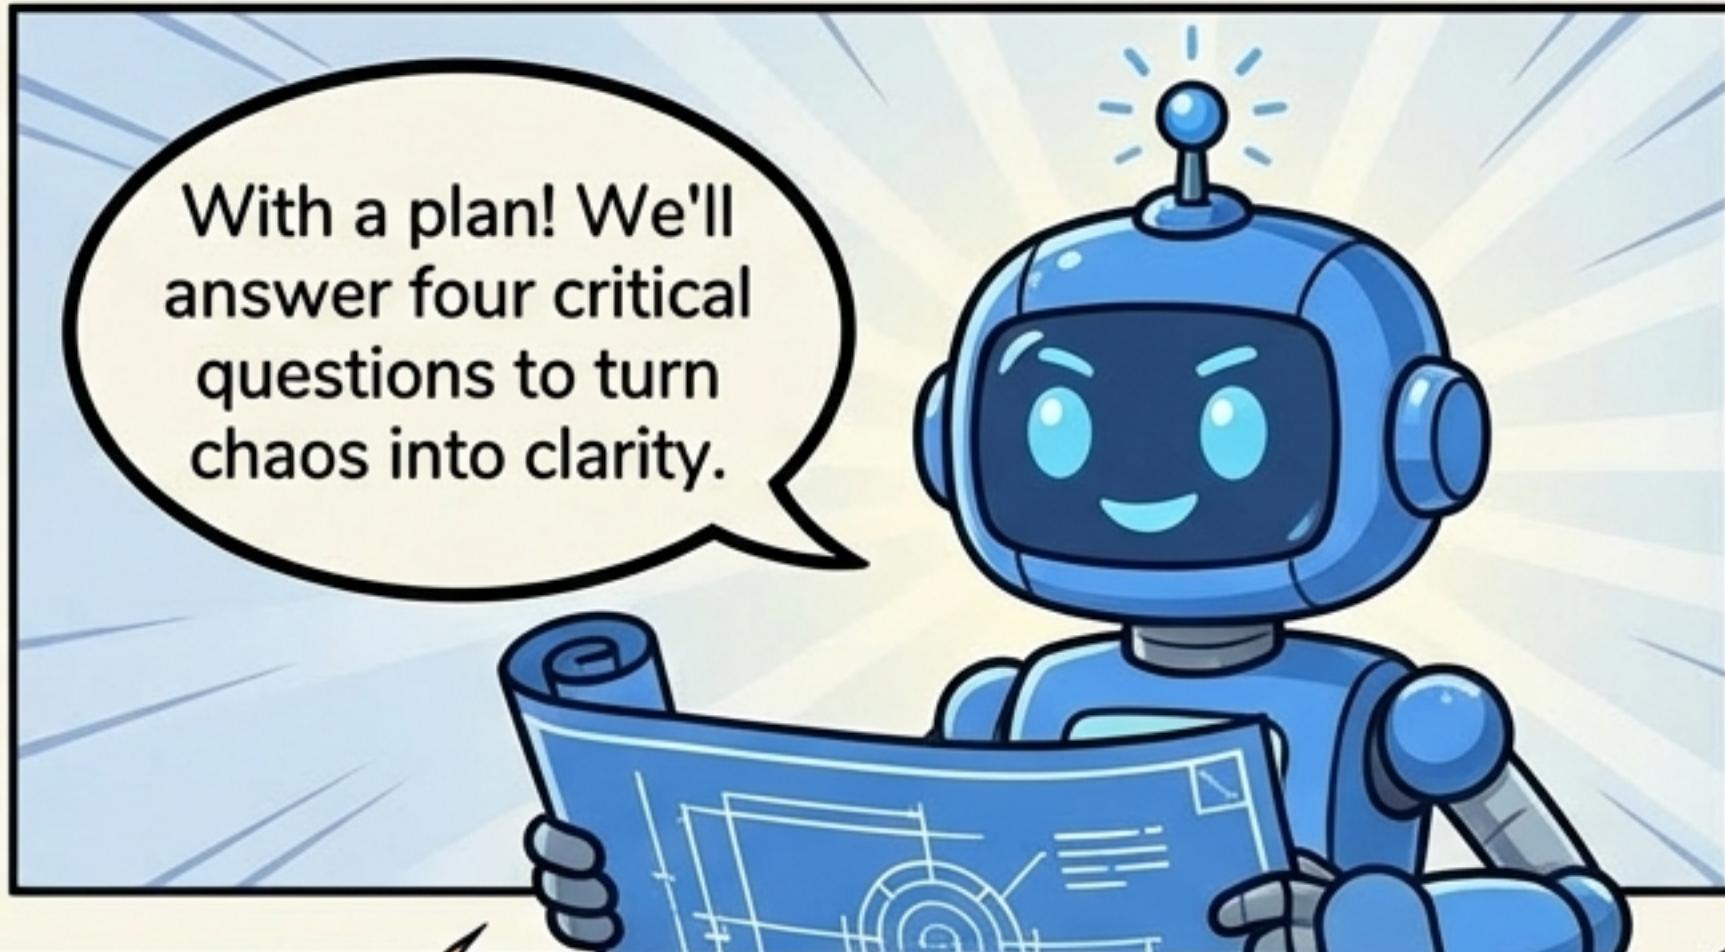


- **Wasted Developer Time:** Endless hours spent diagnosing non-existent code bugs.
- **Delayed Releases:** Critical updates are stalled, waiting for pipelines to pass.
- **Resource Drain:** Multiple job reruns consume valuable machine resources and overload CI servers.

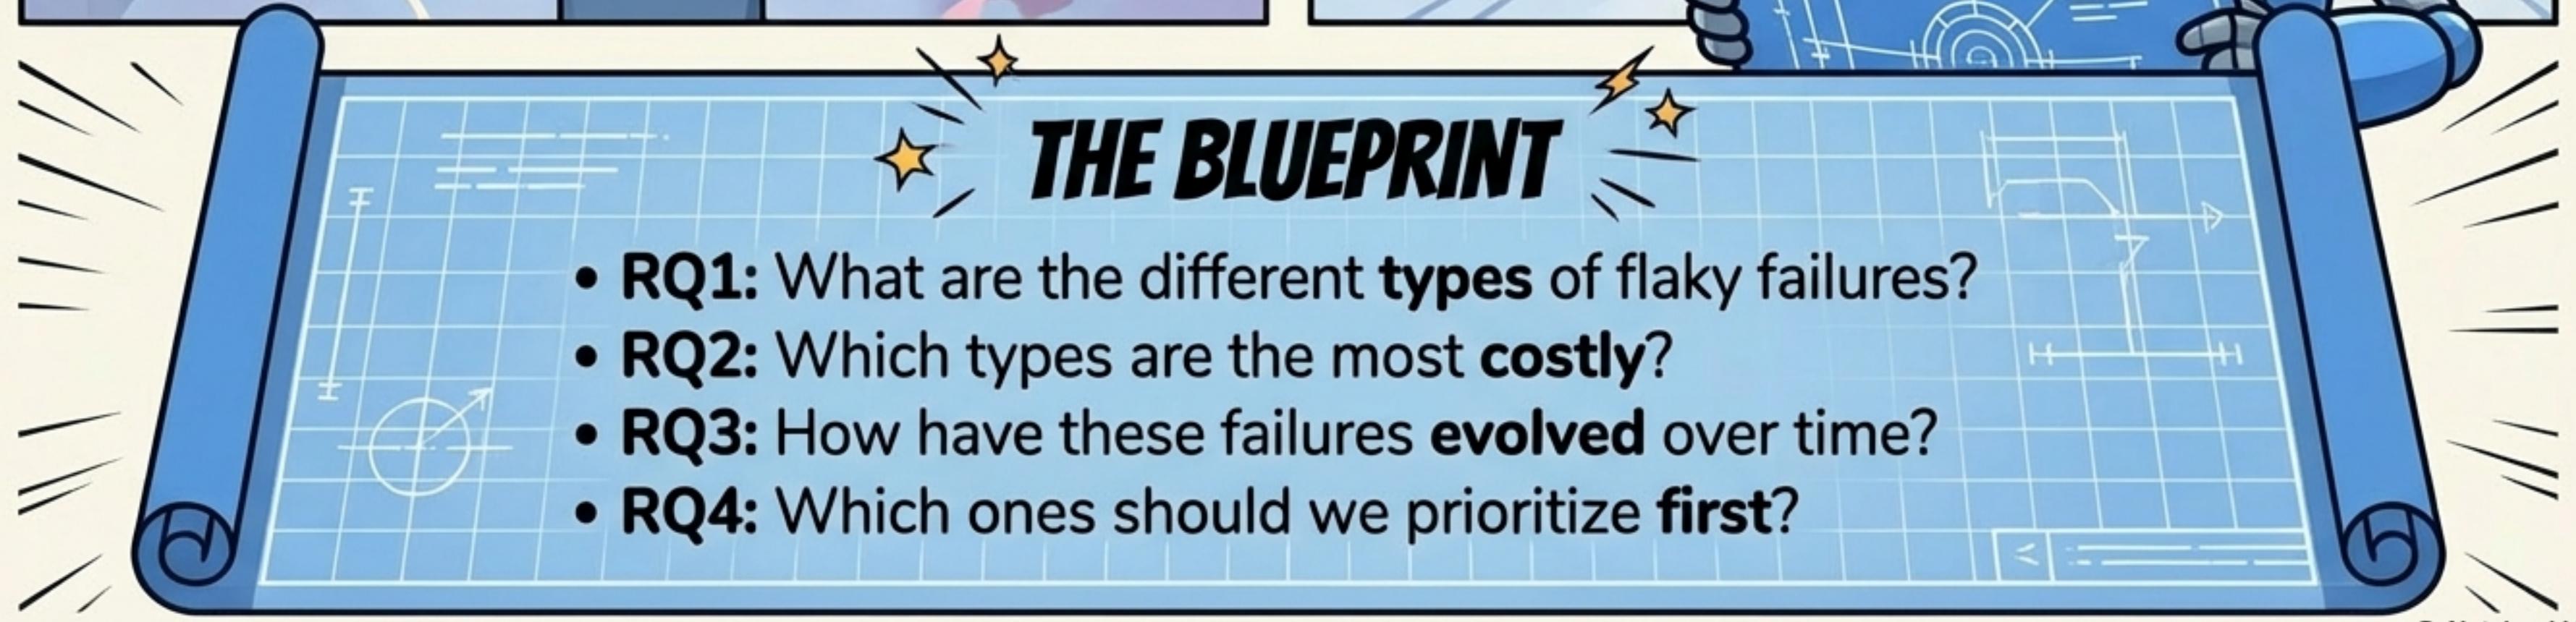
# **TO DEFEAT THE ENEMY, WE MUST FIRST UNDERSTAND IT**



There are so many kinds of flaky failures... Where do we even start?



With a plan! We'll answer four critical questions to turn chaos into clarity.

- 
- ## **THE BLUEPRINT**
- **RQ1:** What are the different **types** of flaky failures?
  - **RQ2:** Which types are the most **costly**?
  - **RQ3:** How have these failures **evolved** over time?
  - **RQ4:** Which ones should we prioritize **first**?

# Our Gadget: The “Flaky-Failure Labeling Gun”

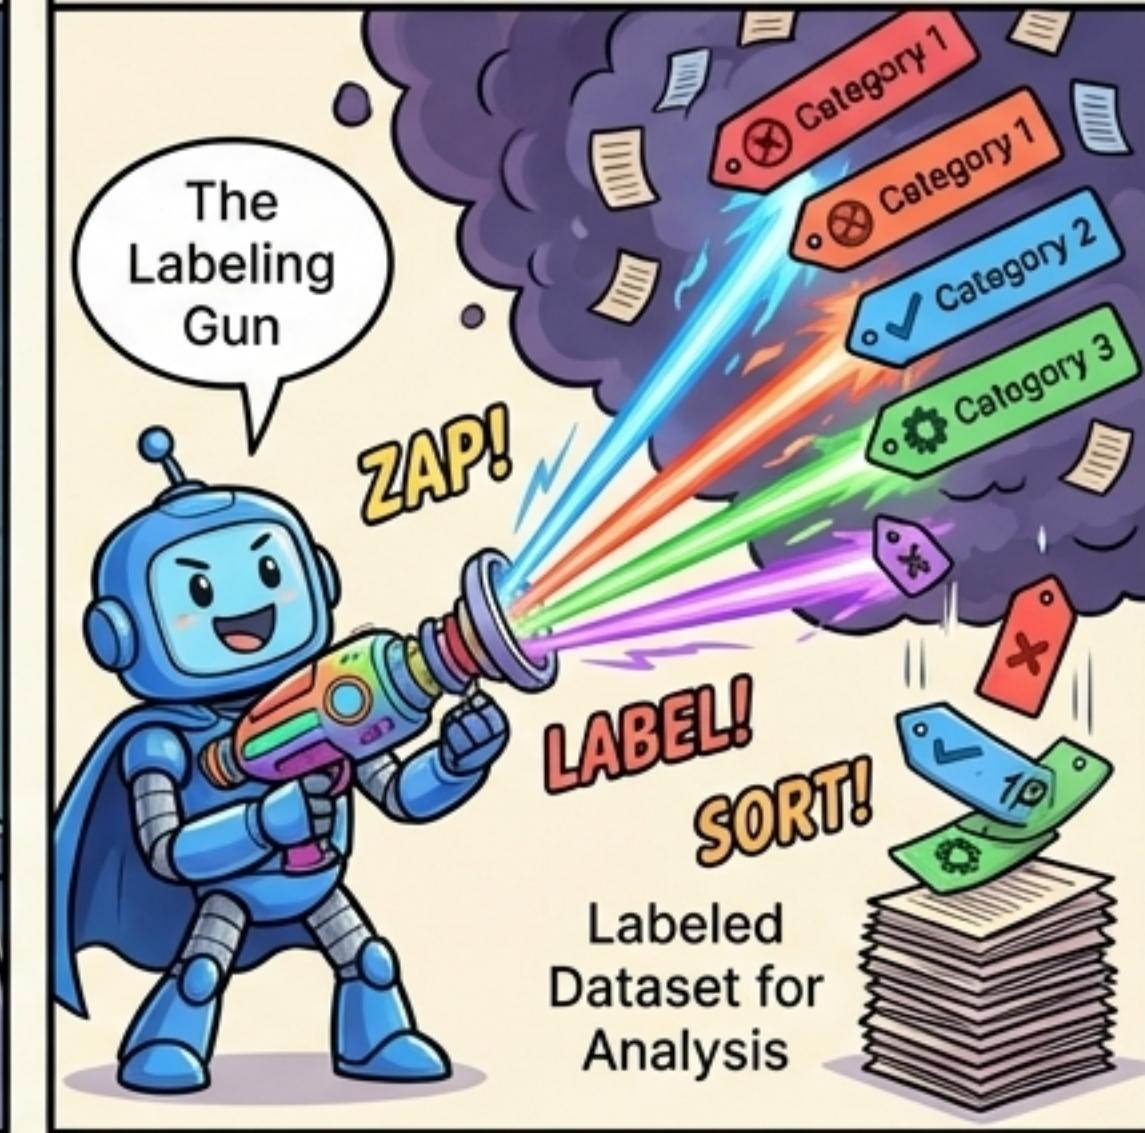
## Collect Samples



## Manual Analysis & Rule Creation



## Automated Labeling

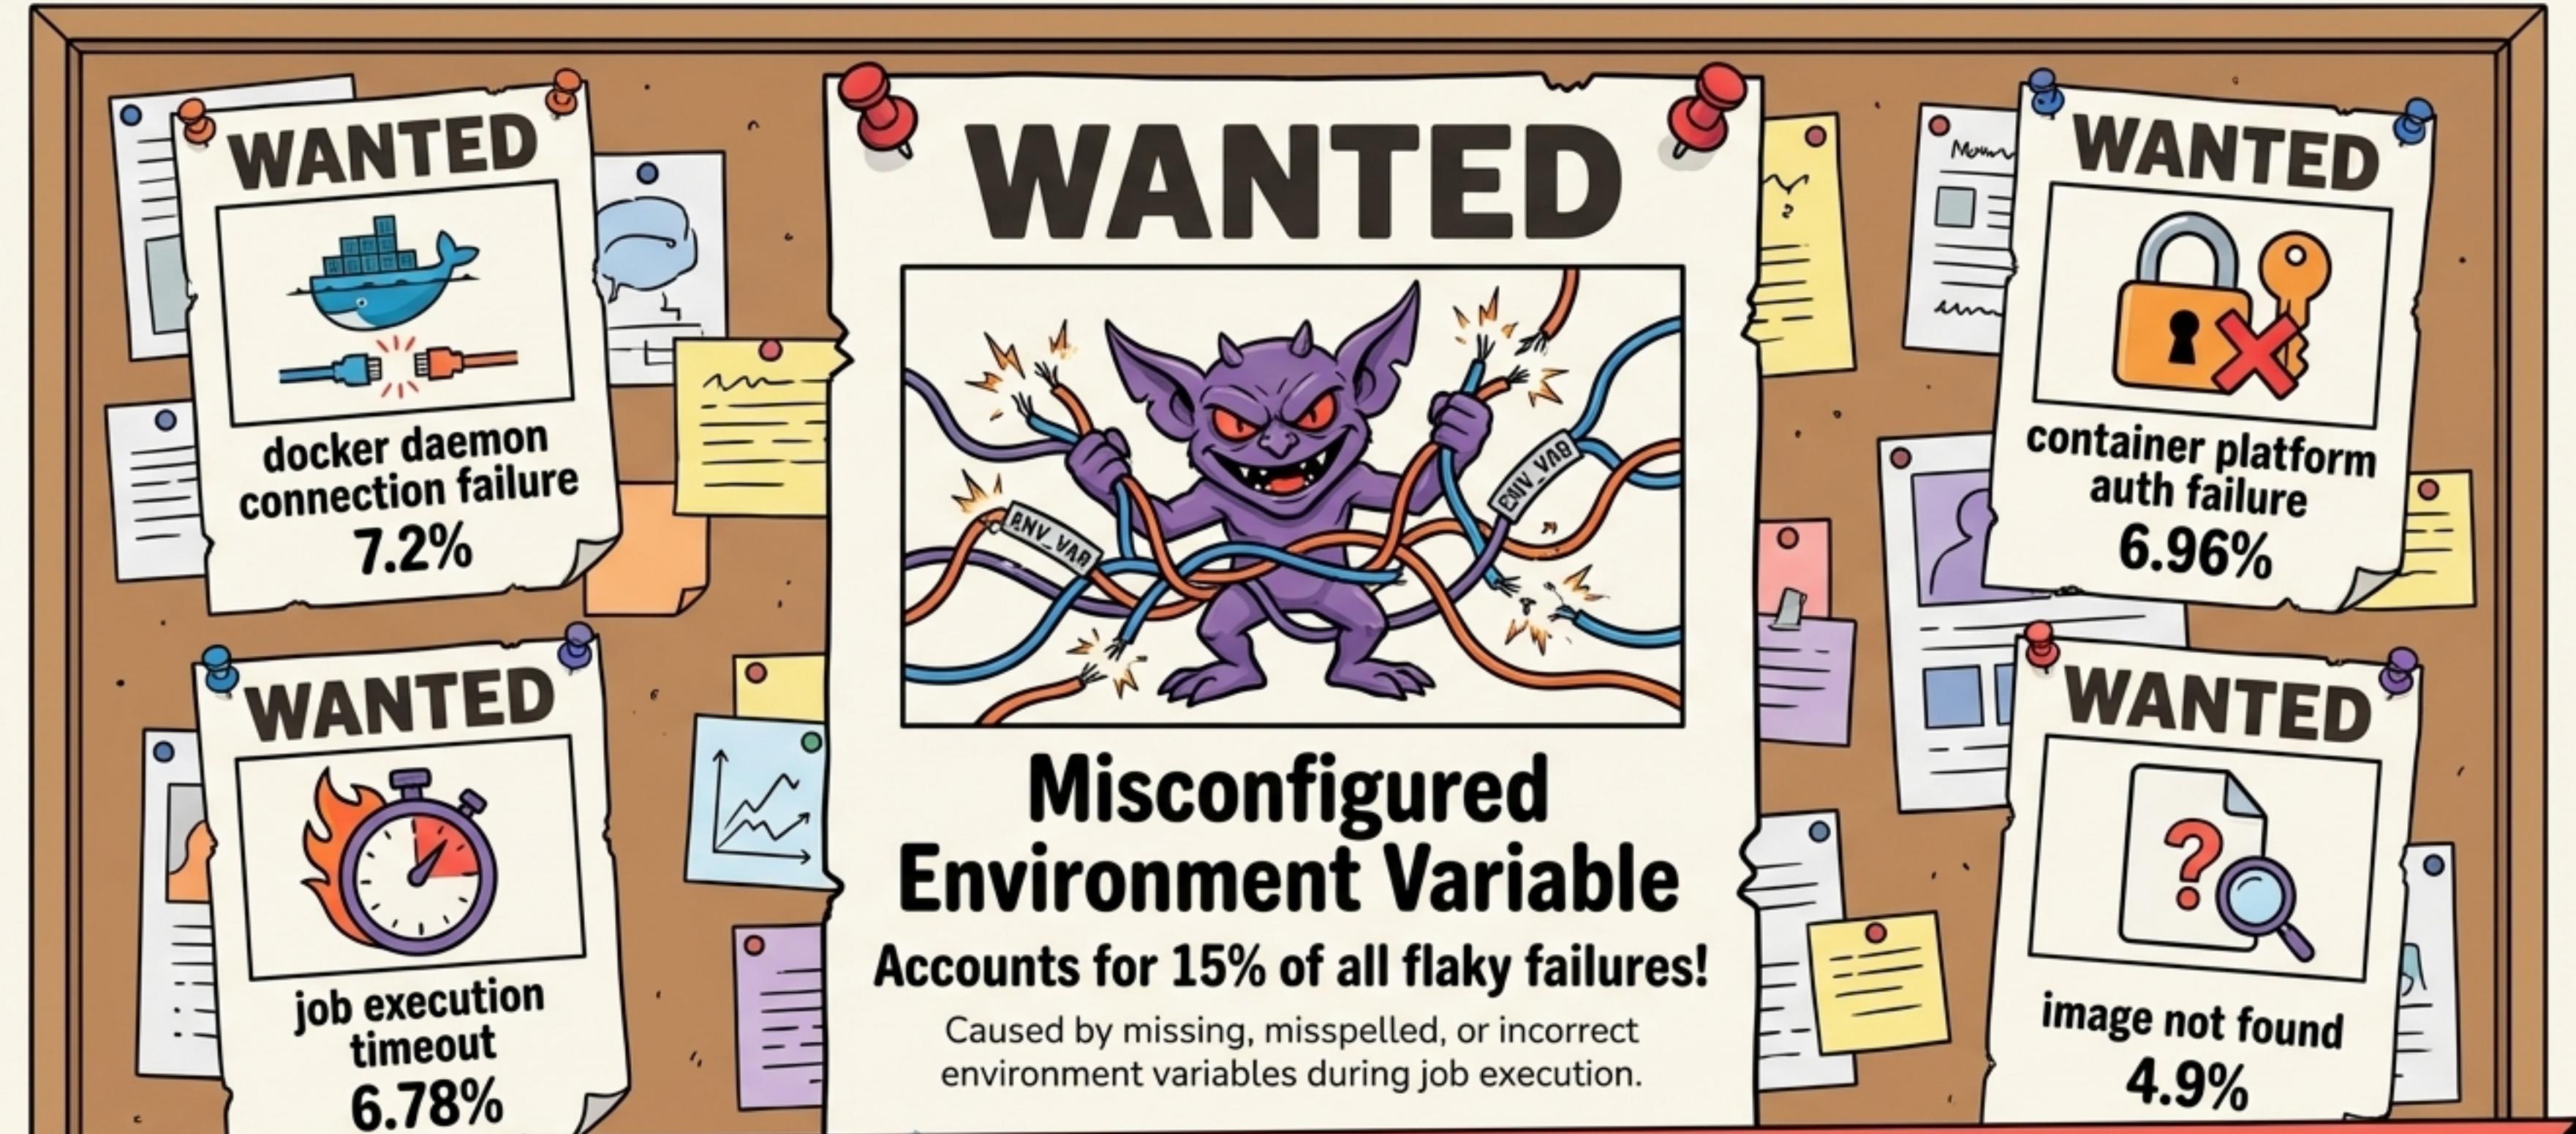


We started with thousands of unlabeled failure logs.

We manually analyzed a statistically significant sample to create 46 distinct categories and regex rules to find them.

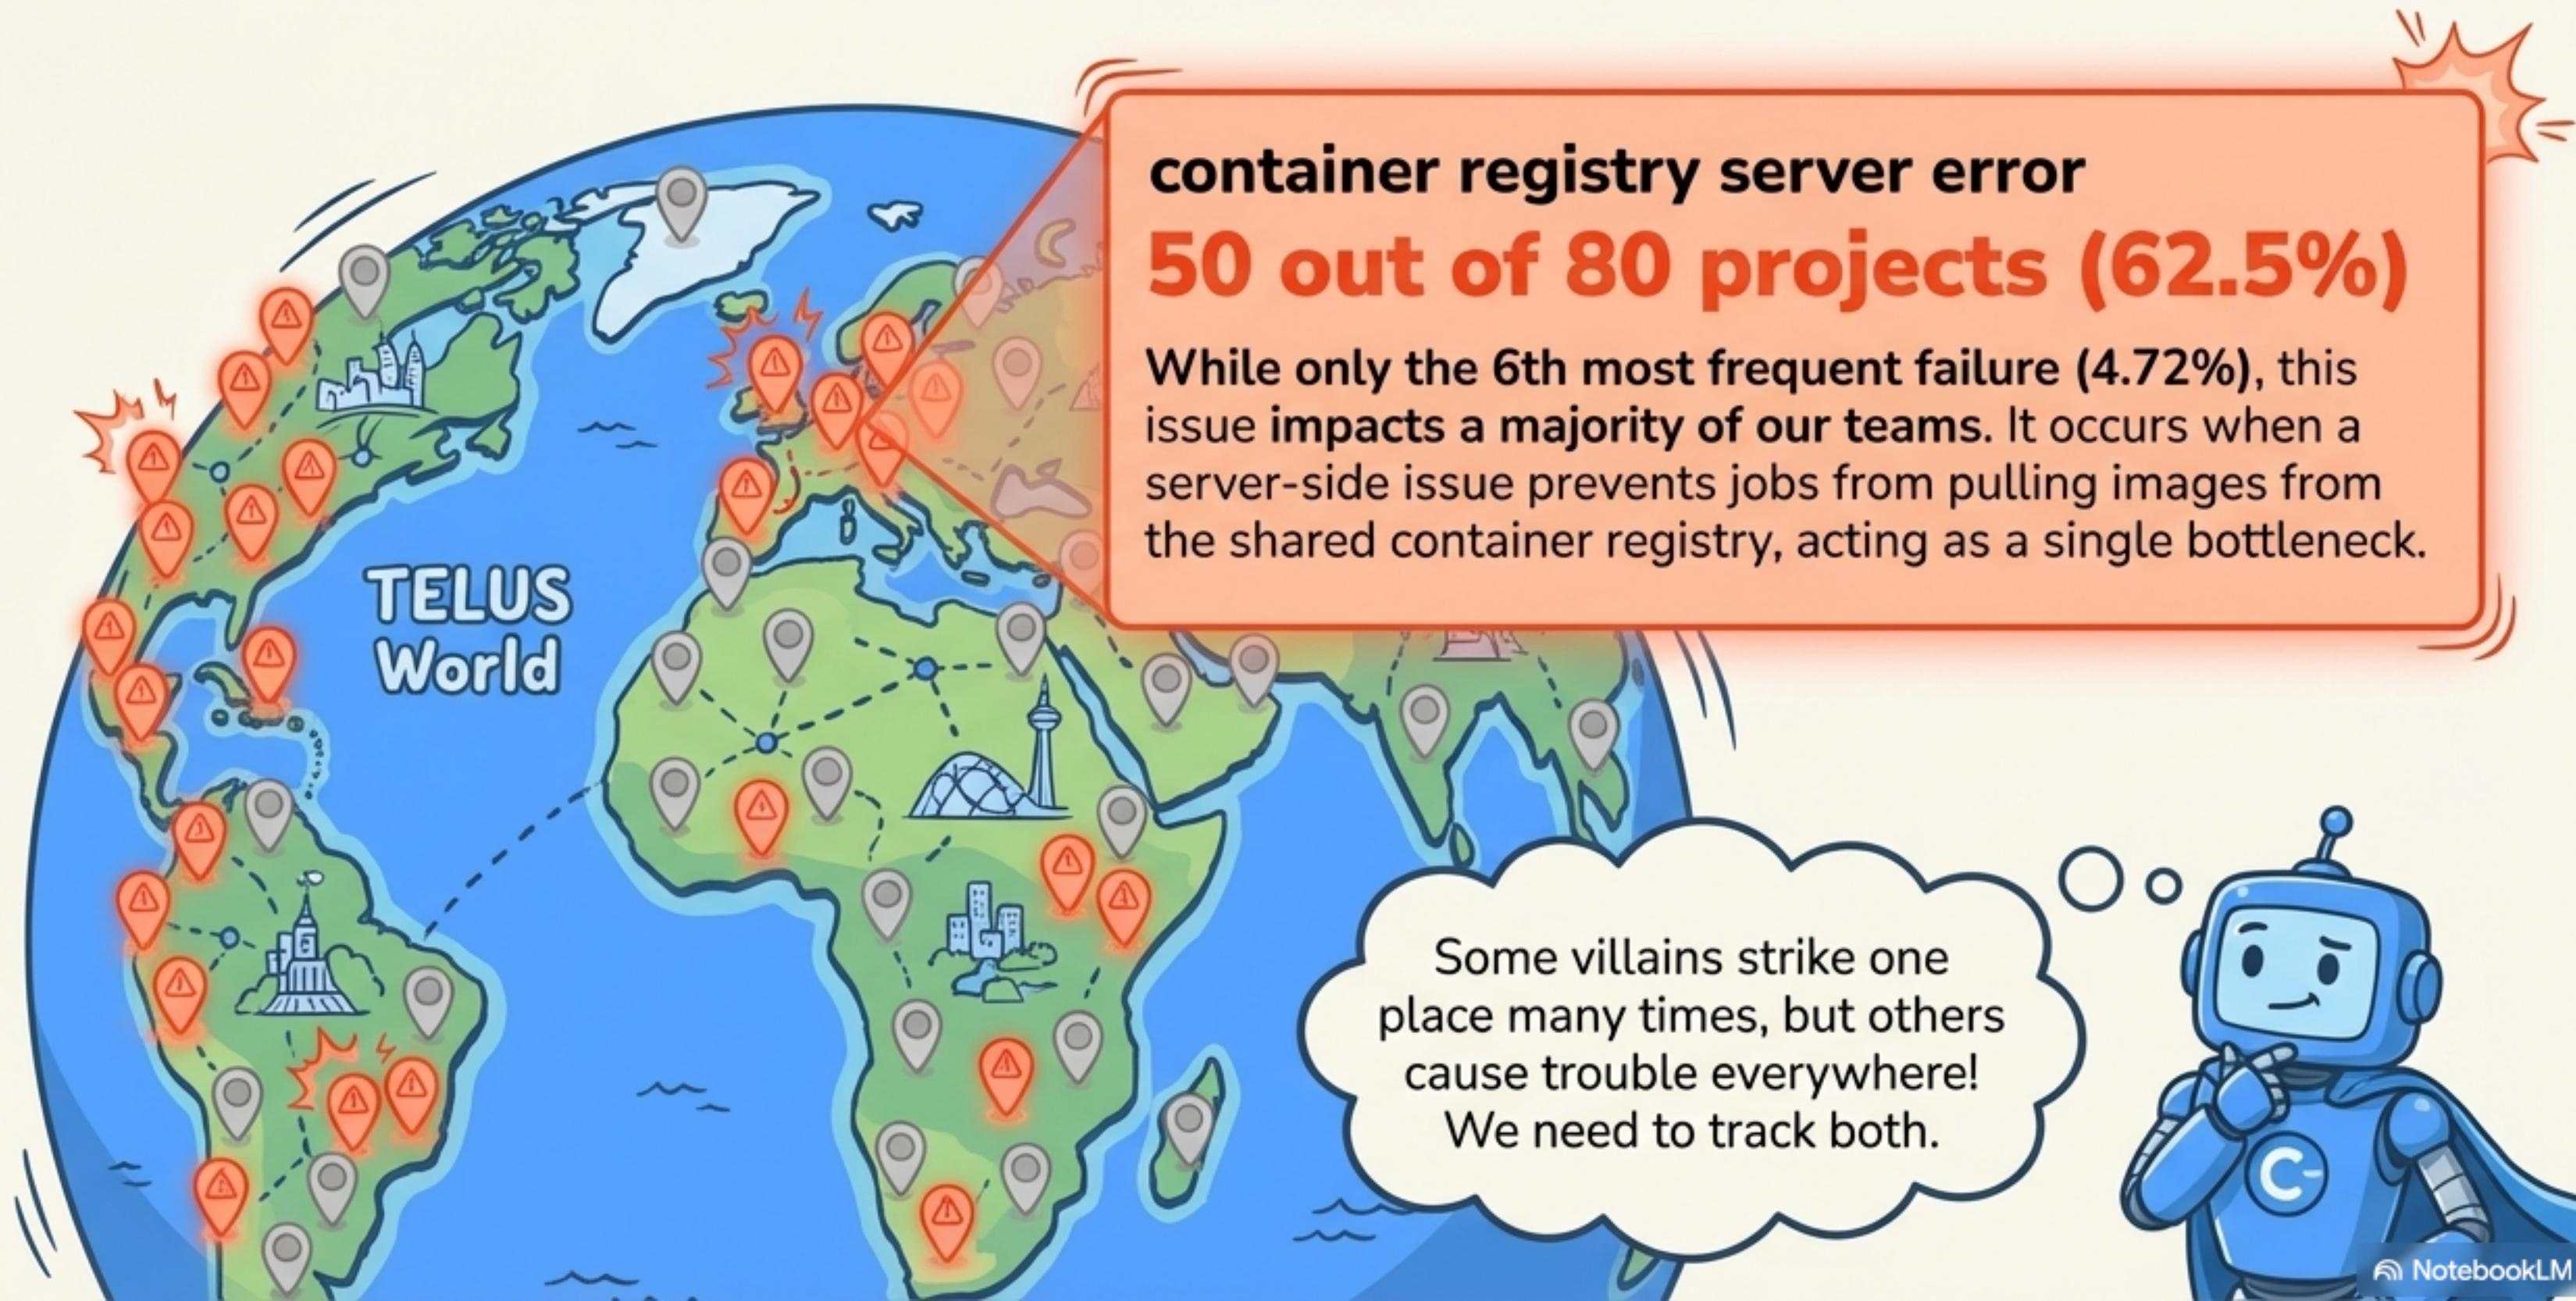
The tool then automatically labeled 4,511 failures with 91% precision, giving us our dataset for analysis.

# The Rogues' Gallery: Meet the 46 Types of Flaky Failures

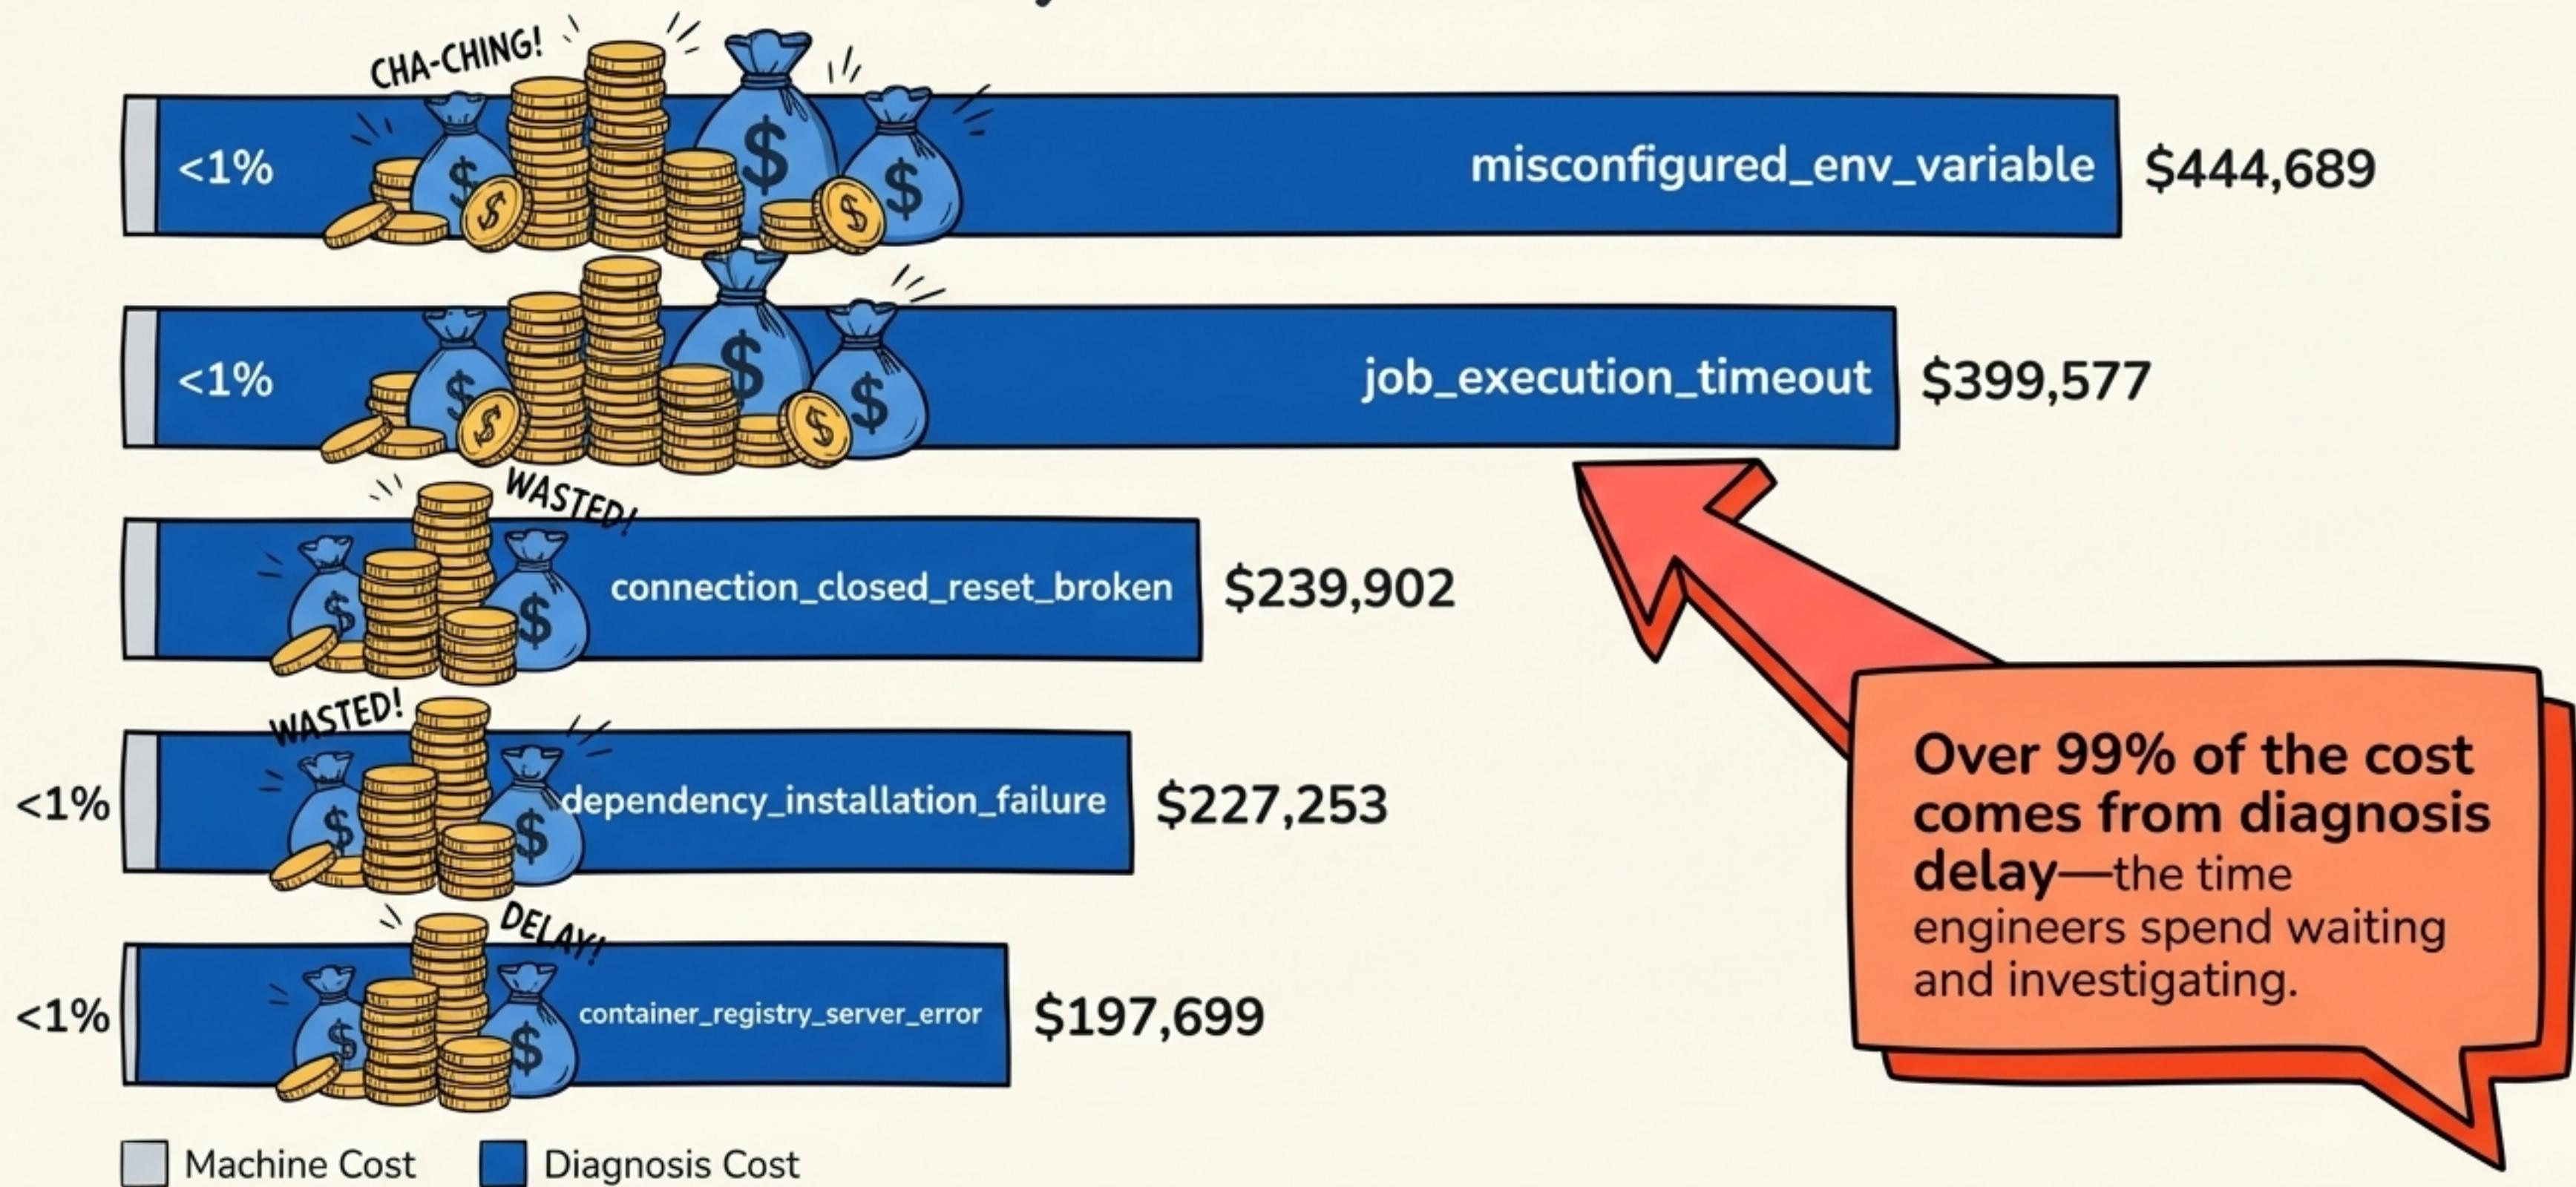


Together, these top 5 culprits are responsible for over 40% of all flaky failures.

# Frequency Isn't Everything: The Most Widespread Gremlin Affects 60% of Projects

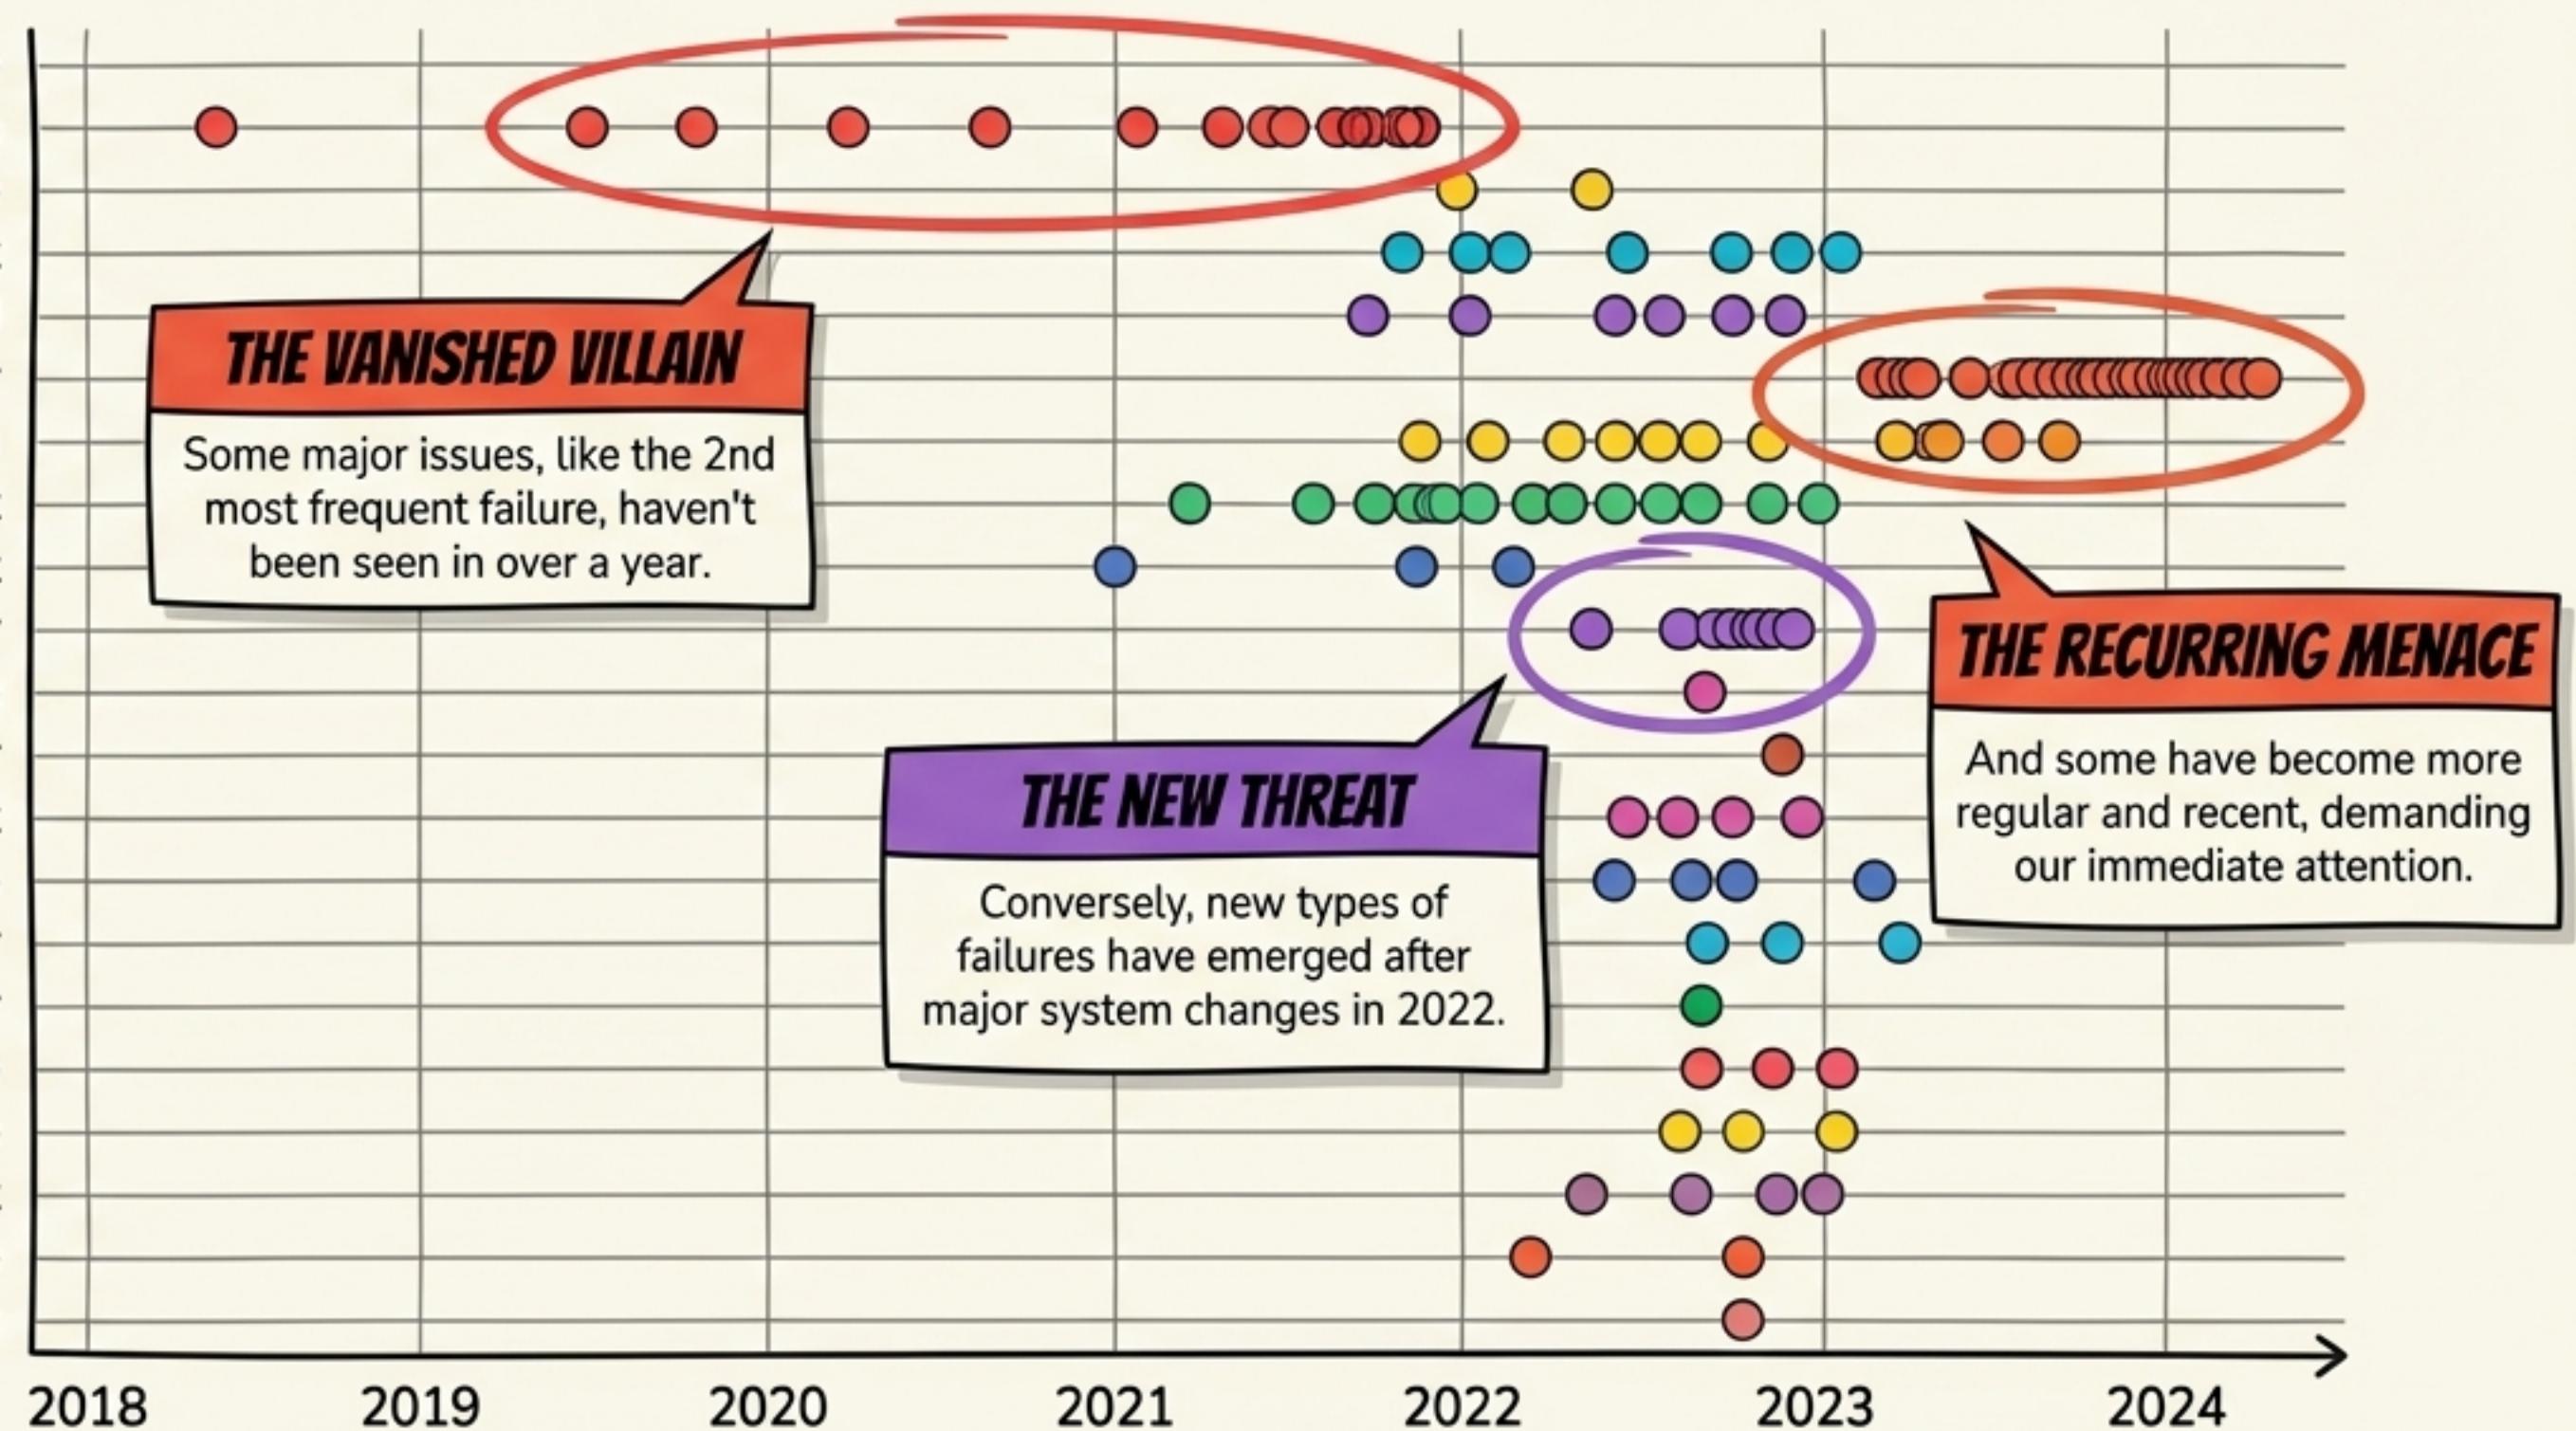


# The True Cost of Flakiness is Measured in Wasted Hours, Not Just Machine Time



# The Battlefield is Always Changing

misconfigured\_env\_variable  
docker\_daemon\_connection\_failure  
container\_platform\_auth\_failure  
job\_execution\_timeout  
image\_not\_found  
container\_registry\_server\_error  
connection\_closed\_reset\_broken  
runner\_pod\_waiting\_timeout  
flaky\_test  
api\_gateway\_deployment\_error  
runner\_pod\_failure  
dependencies\_conflict\_error  
external\_file\_invalid\_format  
dependency\_installation\_failure  
git\_transient\_error  
helm\_resource\_error  
runner\_image\_pull\_failure  
host\_resolution\_failure  
remote\_call\_timeout  
apt\_timezone\_issue  
missing\_logs



# *Time for a New Gadget: The R-F-M Prioritizer!*



## **R - Recency**

How recently have we seen this failure? We focus on the last three occurrences to measure consistent, current threats.



## **F - Frequency**

How often does this failure occur? We count the total number of times it has appeared.



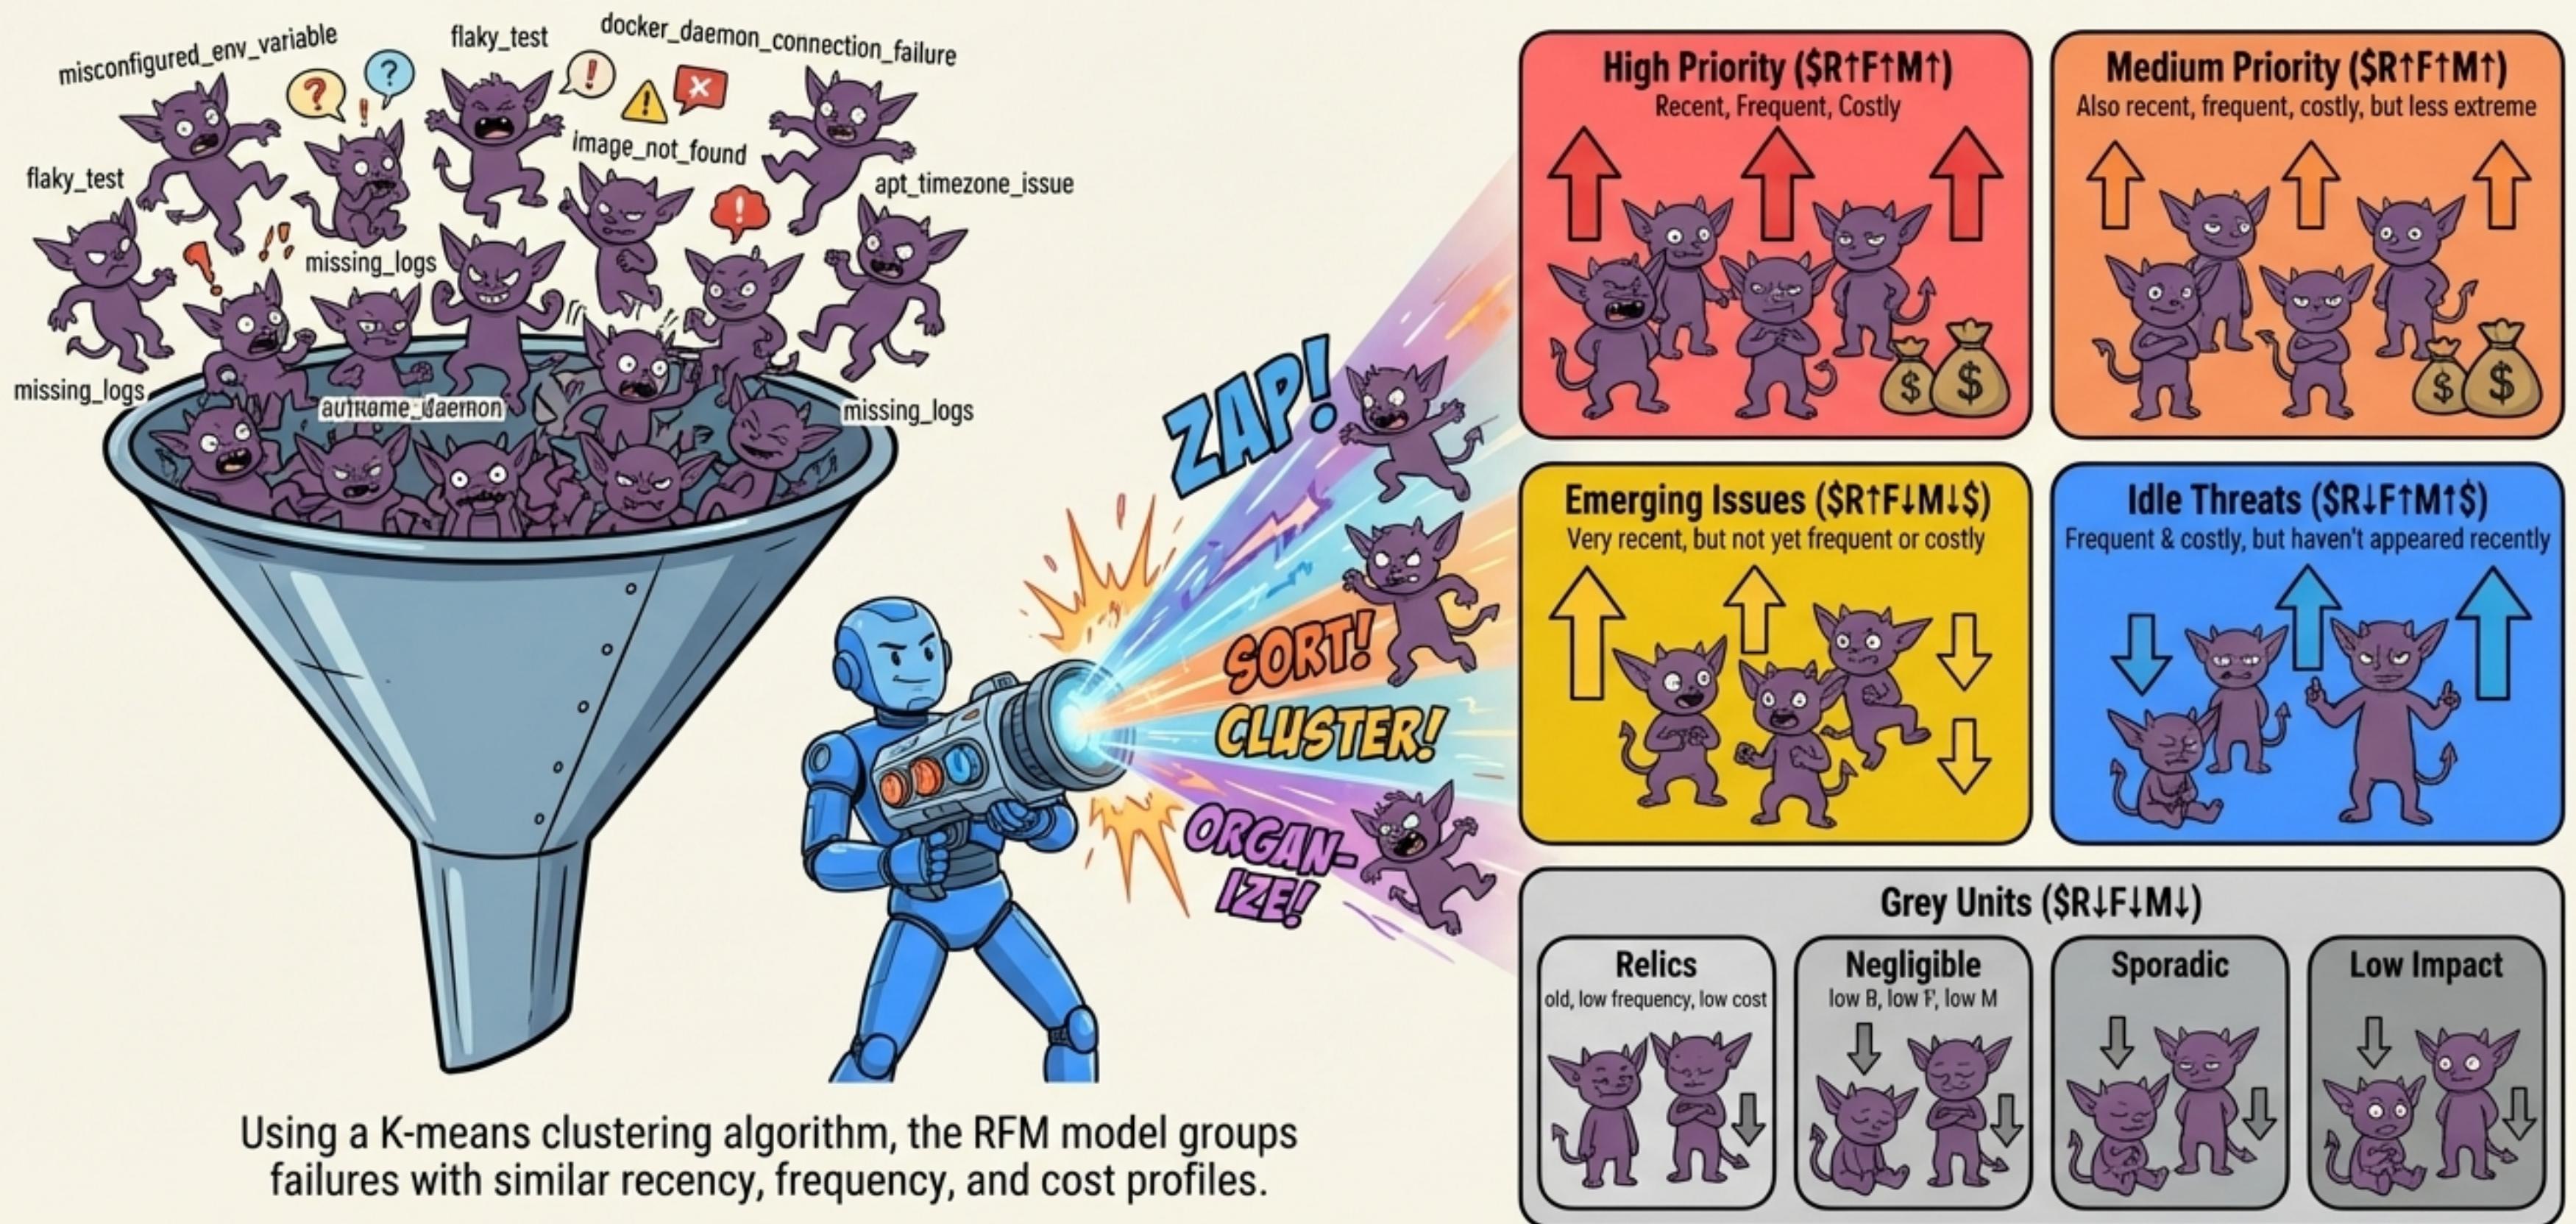
## **M - Monetary**

How costly is this failure? We use our cost model based on diagnosis time and machine resources.



This gadget doesn't just count the gremlins; it tells us which ones are the biggest, most recent, and most expensive threats right now!

# Sorting the Threats: From 46 Categories to 8 Actionable Clusters



# The Battle Plan: Our 14 Highest Priority Targets



TARGET #1: ARCH-VILLAIN

misconfigured  
env variable

R: 27 days | F: 673 | M: \$444k



TARGET #2: ARCH-VILLAIN

job execution  
timeout

R: 8 days | F: 306 | M: \$399k

## HIGH & MEDIUM PRIORITY LIEUTENANTS

- ★ dependency installation failure
- ★ runner pod waiting timeout
- ★ api gateway deployment error
- ★ container registry server error
- ★ flaky test
- ★ git transient error

- ★ external file invalid format
- ★ host resolution failure
- ★ runner image pull failure
- ★ cloud token limit exceeded
- ★ remote call timeout
- ★ helm resource error

Future efforts on automated diagnosis and repair should focus squarely on these 14 categories to maximize impact.

# On the Horizon: 14 More Threats to Keep on Our Radar

## Emerging Issues (\$R↑F↓M↓\$)



These 4 categories (like `runner pod not found`) are very recent. While their frequency and cost are currently low, they represent new problems that could grow over time. We must monitor them closely.

## Idle Threats (\$R↓F↑M↑\$)



These 5 categories (like `runner pod failure`) were highly frequent and costly but haven't surfaced for a while. Monitoring is essential to ensure they are truly resolved and don't re-emerge.

A smart strategy isn't just about fighting today's battles, but also anticipating tomorrow's.

# From Chaos to Clarity: Our Quest is Complete



**BEFORE**



**AFTER**



We Faced the Chaos:  
Acknowledged the 25%  
flaky failure rate.



We Profiled the Enemy:  
Built a taxonomy of 46  
distinct failure types.



We Assessed the Threat:  
Quantified their cost and  
tracked their evolution.



We Forged a Weapon:  
Used the RFM model to  
pinpoint the 14 most critical  
targets.

**We now have a data-driven strategy to systematically reduce waste and improve the reliability of our CI/CD pipelines.**

# The Blueprint for Conquering Flaky Failures

## A Taxonomy for Diagnosis



We provide a comprehensive taxonomy of 46 flaky failure categories, offering a starting point for any organization to begin diagnosing their own CI/CD issues.

## The RFM Prioritization Framework



We introduce a novel approach using Recency, Frequency, and Monetary (RFM) analysis to prioritize technical issues, a framework that can be adapted to many other contexts beyond flaky jobs.

This work lays the foundation for building targeted, automated solutions that fix the problems that matter most.